

Light

@rantav
totango



Spotify™

WHO AM I?

An aerial photograph of a coastal landscape during sunset or sunrise. The sky is filled with soft, pastel-colored clouds, transitioning from blue to orange and yellow near the horizon. In the foreground, there are dark, rugged mountains and hills. A large body of water, possibly a lake or a bay, reflects the warm colors of the sky. The overall atmosphere is serene and contemplative.



WHO AM I?

- A developer
 - Google, Microsoft, Outbrain, Gigaspaces, Totango etc
 - Hector, flask-restful-swagger, meteor-migrations, monitoring...
- Podcast: reversim.com
- devdev.io
- Gormim

A nighttime satellite view of Earth from space, showing city lights and clouds.

WHAT IS LUIGI?

WHAT IS LUIGI?



WHAT IS LUIGI?

- A **Workflow Engine**.
 - Who the fuck needs a workflow engine?



WHAT IS LUIGI?

- A **Workflow Engine**.
 - Who the fuck needs a workflow engine?
 - You do!!!
 - If you run hadoop (or other ETL jobs)
 - If you have dependencies b/w them (who doesn't?!)
 - If they fail (`s/if/when/`)



WHAT IS LUIGI?

- A **Workflow Engine**.
 - Who the fuck needs a workflow engine?
 - You do!!!
 - If you run hadoop (or other ETL jobs)
 - If you have dependencies b/w them (who doesn't?!)
 - If they fail (**s/if/when/**)
 - Luigi doesn't replace Hadoop, Scalding, Pig, Hive, Redshift.
 - It **orchestrates** them



SCREENSHOTS

Luigi Task Visualiser

<https://pipeline.spotify.net/luigi/static/visualiser/index.html>

Failed Tasks

- WebplayerRequestTimeImporterMulti (1)
- WeeklyBrowseData (1)

Running Tasks

- ActivationDomino (1)
- AggregateClientCrashHockeyApp (10)
- CrashLogsToHdfs (15)
- DominolImportTask (1)
- ImpressionsAndUsers (0)
- OrganicMultiplier (2)
- RetentionDomino (10)
- UpdateLogAvroSchemas (0)
- UserUsageSummary (1)
- UTMAcquisitionShard (0)

AggregateClientCrashHockeyApp (test=False, date=2013-11-02) 6:32:03 PM

AggregateWebPlayerUpgradeTest (1)

AndroidHockeyAppCrashes (10)

BandwidthConnectionJoin (1)

DominolImportTask (2)

DumpFslImage (1)

DumpMetadata (2)

DumpMetadata (date=2013-11-06, table=publishedaudio) 7:21:09 PM

DumpMetadata (date=2013-11-06, table=audio) 7:21:09 PM

DumpSuperviseStatus (0)

Luigi Task Status Active tasks

[https://pipeline.spotify.net/luigi/static/visualiser/index.html#HockeyAppCrashesToPostgresMulti\(last_date=2013-11-06, days_back=15, days_back_strict=10, today=False\)](https://pipeline.spotify.net/luigi/static/visualiser/index.html#HockeyAppCrashesToPostgresMulti(last_date=2013-11-06, days_back=15, days_back_strict=10, today=False))

Dependency Graph

Legend: Failed (Red), Pending (Yellow), Running (Blue), Done (Green)

HOW DO YOU ETL YOUR DATA?

- Hadoop
- Spark
- Redshift
- Postgres
- Ad-hoc java/python/ruby/go/...

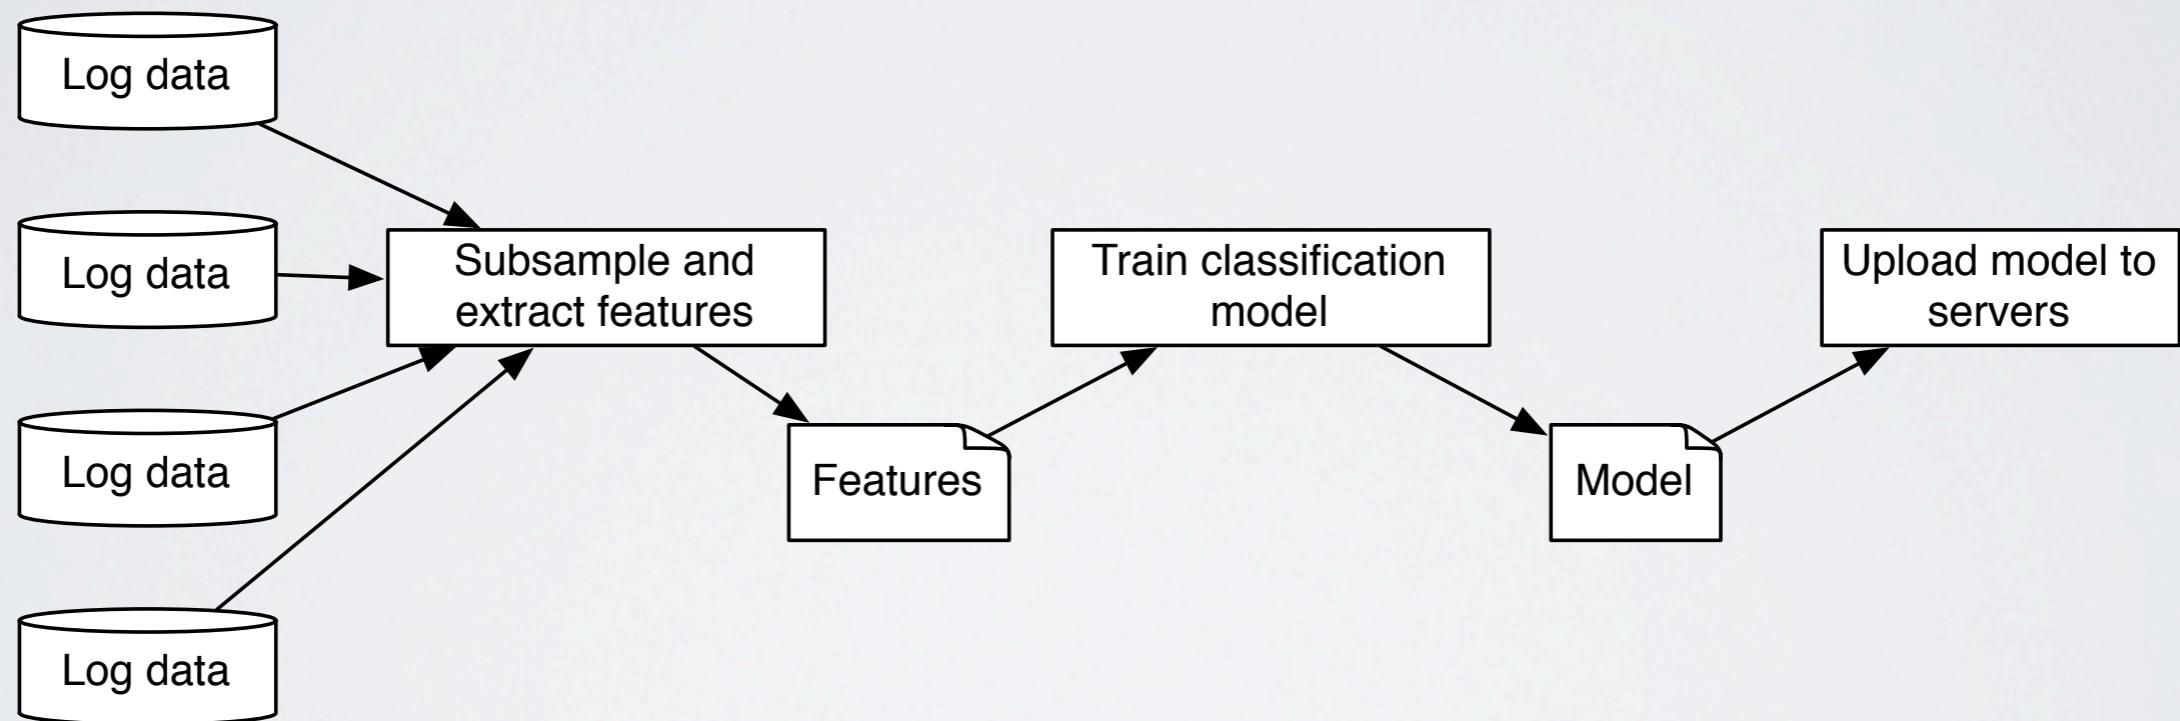
RUNNING **ONE** JOB IS EASY RUNNING **MANY** IS HARD

- 100s of concurrent jobs, 1000s Daily.
- Job dependencies.
 - E.g. first copy the file, then crunch it.
- Errors / retries
 - Idempotency
- Monitoring / Visuals

THE WRONG WAY TO DO IT



EXAMPLE WORKFLOW



THE CRON PHENOMENON

THE **WRONG** WAY TO DO IT

```
# m h dom mon dow user  command
0 1 * * * spotify-analytics /usr/bin/spotify_extract_features --date `date +%Y-%m-%d` # let's hope the log files are available by 1am
0 2* * * spotify-analytics /usr/bin/spotify_train_model --date `date +%Y-%m-%d` # hopefully previous job finished by 2am
0 3 * * * spotify-analytics /usr/bin/spotify_update_model --date `date +%Y-%m-%d` # same here... OMG ugly!!!

```

-uu:---F1 crontab All L3 (Fundamental)-----
Wrote /Users/erikbern/pydata_pres/crontab

THE CRON PHENOMENON

THE **WRONG** WAY TO DO IT

```
# m h dom mon dow user  command
0 1 * * * spotify-analytics /usr/bin/spotify_extract_features --date `date +%Y-%m-%d` # let's log files are available by 1am
0 2 * * * spotify-analytics /usr/bin/spotify_train_model --date `date +%Y-%m-%d` # this job finished by 2am
0 3 * * * spotify-analytics /usr/bin/spotify_update_model --date `date +%Y-%m-%d` # ugly!!!
uu:---F1 crontab      All L3  (Fundamental)-----
Wrote /Users/erikbern/pydata_pres/crontab
```

Don't try this at home!!!

ENTER LUIGI



ENTER LUIGI

- Like Makefile - but in python
 - And - For data
 - Integrates well with data targets
 - Hadoop, Spark, Databases
 - Atomic file/db operations
 - Visualization
 - CLI - really nice developer interface!



LUIGITASK

```
class MyTask(Task):
    def output(self):
        pass

    def requires(self):
        pass

    def run(self):
        pass
```

LUIGITASK

```
class AggregateArtists(Task):
    def output(self):
        return HdfsTarget("data/artist_streams.tsv")
    def requires(self):
        return Streams()
    def run(self):
        with self.input().open('r') as in_file:
            ... # read stuff from in_file
        with self.output().open('w') as out_file:
            ... # write stuff to out_file
    if __name__ == "__main__":
        luigi.run()
```

Store output in HDFS

Run - actual work

Read from input

Write to output

Run with __main__

RUN FROM THE CLI

Run on the command line:

```
$ python dataflow.py Aggreg
```

```
DEBUG: Checking if AggregateArtists() is complete
INFO: Scheduled AggregateArtists()
DEBUG: Checking if Streams() is complete
INFO: Done scheduling tasks
DEBUG: Asking scheduler for work...
DEBUG: Pending tasks: 1
INFO: [pid 74375] Running      AggregateArtists()
INFO: [pid 74375] Done          AggregateArtists()
DEBUG: Asking scheduler for work...
INFO: Done
INFO: There are no more tasks to run at this time
```

TASK PARAMETERS

```
class AggregateArtists(Task):  
    date = DateParameter()
```

```
$ python dataflow.py AggregateArtists --date 2013-03-05
```

AWESOME HADOOP (MR) SUPPORT

```
class AggregateArtists(luigi.hadoop.JobTask):
    def requires(self):
        return Streams()

    def output(self):
        return HdfsTarget("data/artist_streams.tsv")

    def mapper(self, line):
        timestamp, artist, track = line.split('\t')
        yield artist, 1

    def reducer(self, artist, streams):
        yield artist, sum(streams)
```

WEB UI

Luigi Task Status

Task List Dependency Graph

Upstream Failure

Failed Tasks

- + 1 AccountIndexDataLoaderTask
- + 3 ActiveListRealtimeTask
- + 1 AvgActivityLevelTask
- + 2 CreateUserDocumentTask
- + 1 DownloadSdrTask
- + 1 HadoopAccountAggregationTask
- + 1 HadoopUserAggregationTask
- + 2 UserDocumentsMergeTask

Pending Tasks

- + 1 AccountAgingTask
- + 3 AccountIndexDataLoaderTask
- + 1 AccountSecondIndexDataLoaderTask

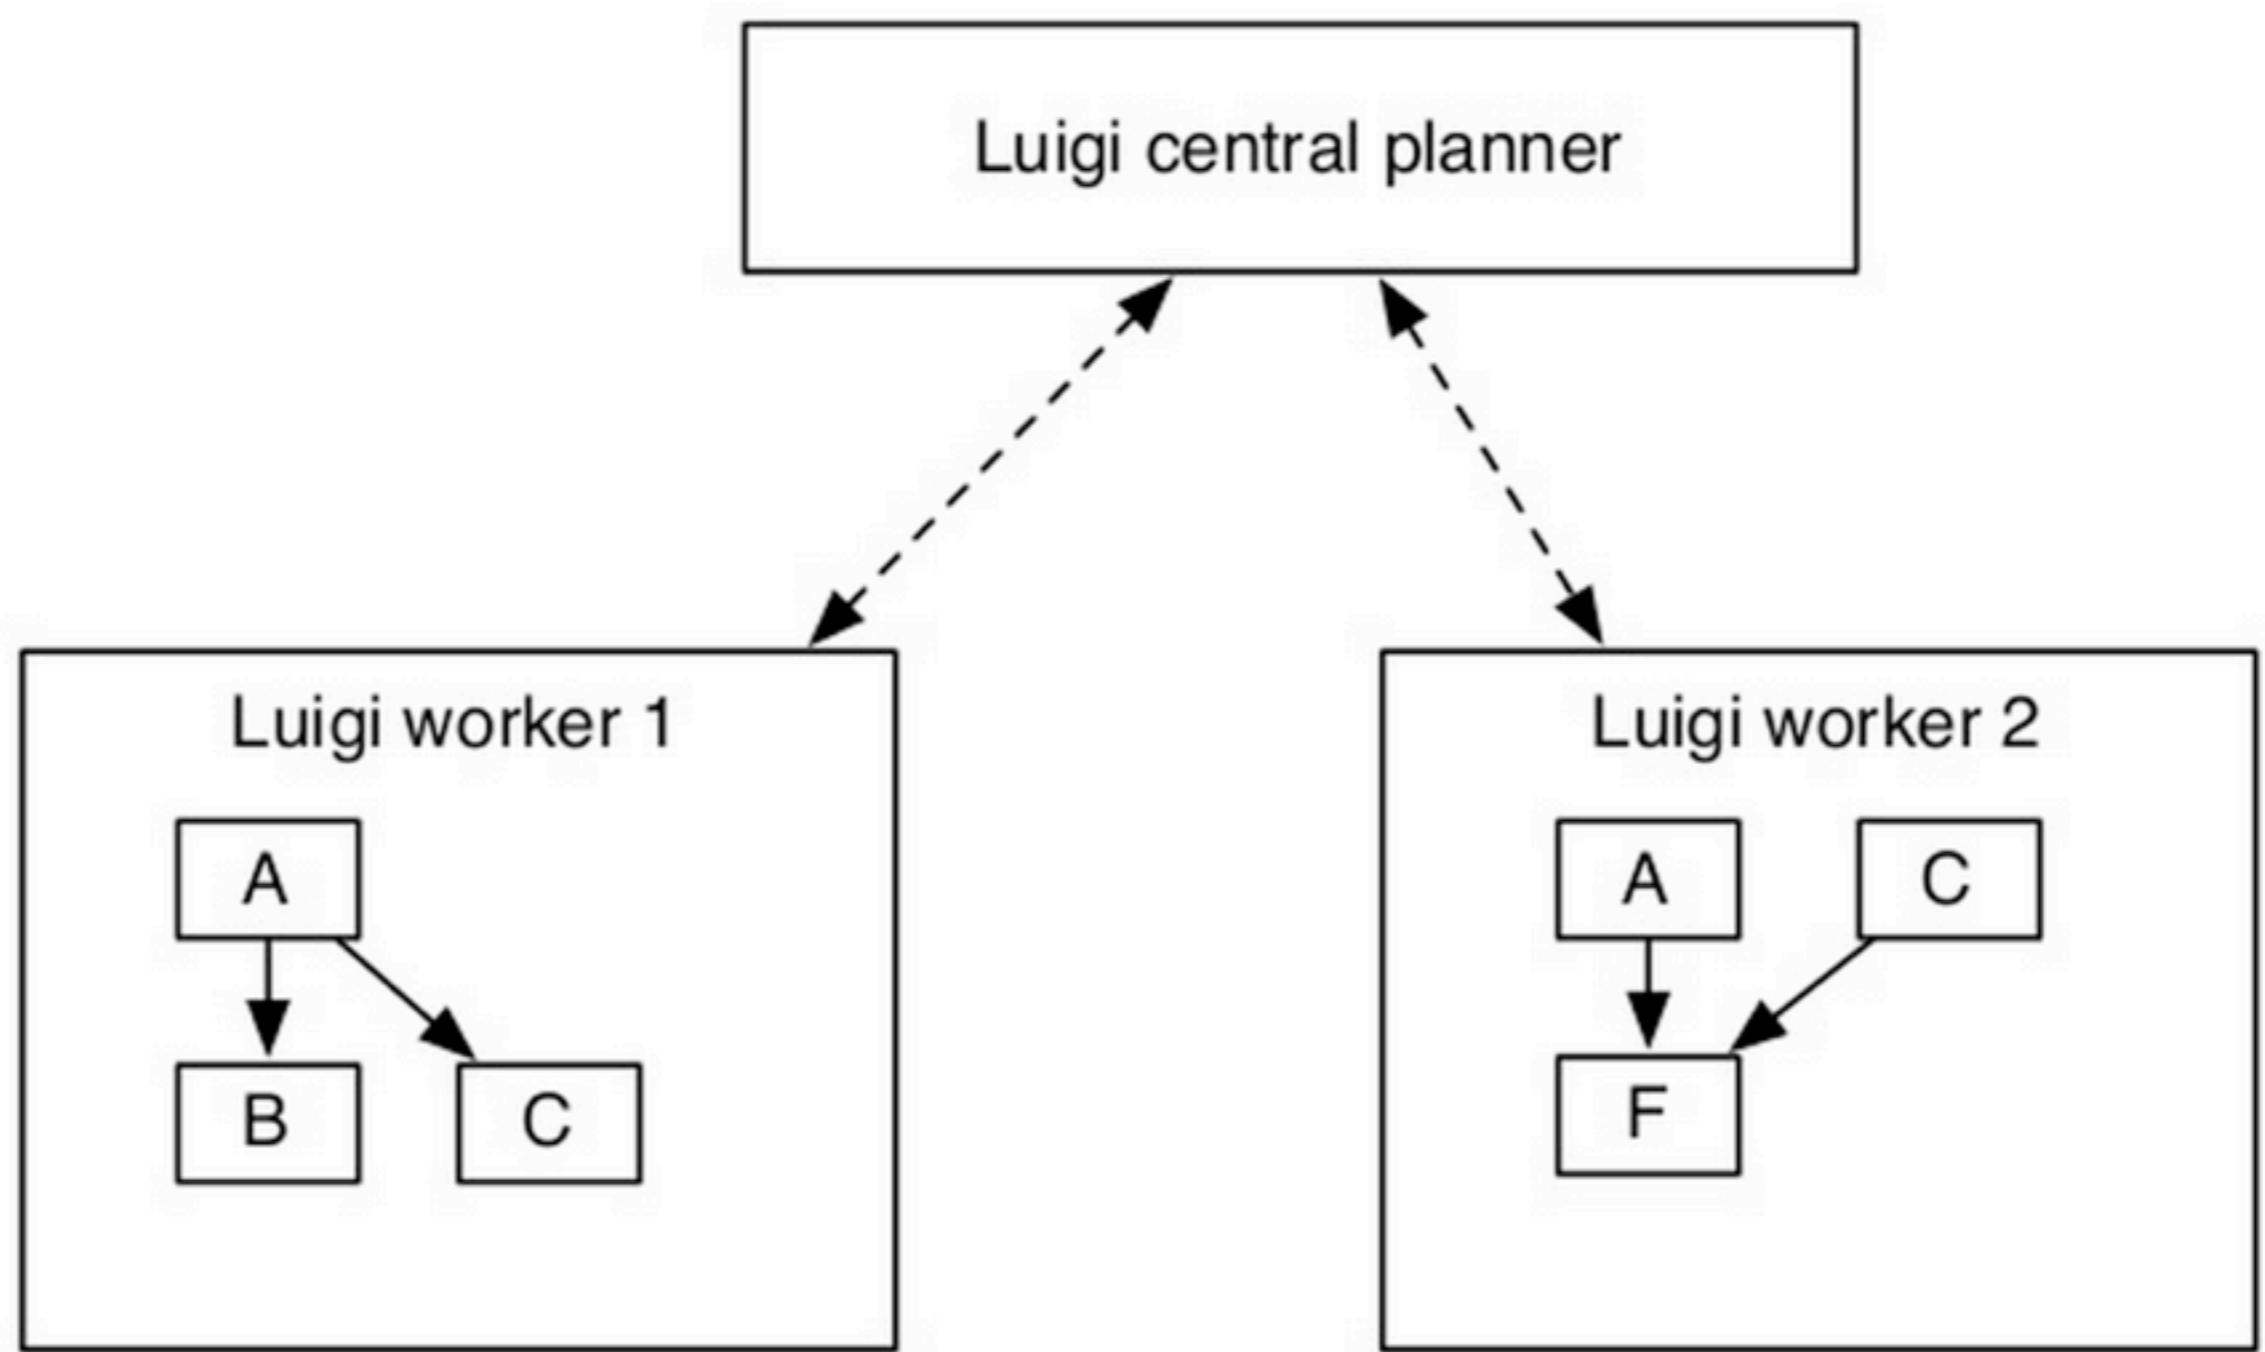
AllDailyTask(service_id=..., day=2014-06-21, force_run=False, tz_string=-0600)

Dependency Graph

Legend:
● Failed
● Running
● Pending
● Done

```
graph TD; AllDailyTask --> ActiveListOrganizationStateTask; AllDailyTask --> ClearRealtimeAggregationsTablesTask; AllDailyTask --> CreateAccountDocumentTask; AllDailyTask --> PeriodicEndMarkerCommand; AllDailyTask --> AddCompressorTask; AllDailyTask --> TotangoAttributesTask; AllDailyTask --> UpdateOrganizationSessionTask; AllDailyTask --> UserEngagementScoreTask; AllDailyTask --> DailyUserDocumentsIndexTask; AllDailyTask --> HistoricalAccountIndexTask; AllDailyTask --> UserDocumentsMergeTask;
```

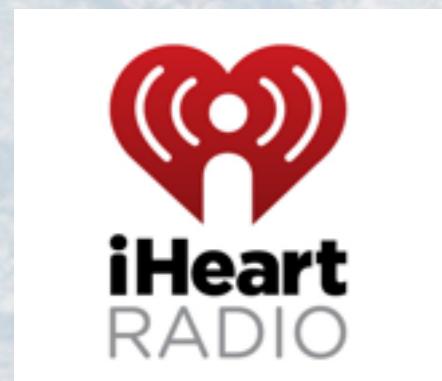
PROCESS SYNCHRONIZATION



USED BY



Spotify®



foursquare®



SEMI-DEEP DIVE

A photograph of a scuba diver swimming in the ocean. The diver is wearing a dark wetsuit and fins, and is carrying a tank and regulator. A large cloud of bubbles is visible behind them, indicating their recent movement. The water is a deep blue-green color.

Programming for Luigi

LUIGI TASKS

- Implement 4 method:

def input(self) (optional)

def output(self)

def run(self)

def depends(self)

LUIGI TASKS

- Or extend one of the predefined tasks
 - **S3CopyToTable**
 - **RedshiftManifestTask**
 - **SparkJob**
 - **HiveQueryTask**
 - **HadoopJobTask**

EXAMPLE LOCAL WORDCOUNT

```
class WordCount(luigi.Task):
    date_interval = luigi.DateIntervalParameter()

    def requires(self):
        return [InputText(date) for date in self.date_interval.dates()]

    def output(self):
        return luigi.LocalTarget('/var/tmp/text-count/%s' % self.date_interval)

    def run(self):
        count = {}
        for file in self.input():
            for line in file.open('r'):
                for word in line.strip().split():
                    count[word] = count.get(word, 0) + 1

        # output data
        f = self.output().open('w')
        for word, count in count.iteritems():
            f.write("%s\t%d\n" % (word, count))
        f.close()
```

EXAMPLE

HADOOP WORDCOUNT

```
class WordCount(luigi.hadoop.JobTask):
    date_interval = luigi.DateIntervalParameter()

    def requires(self):
        return [InputText(date) for date in self.date_interval.dates()]

    def output(self):
        return luigi.hdfs.HdfsTarget('/tmp/text-count/%s' % self.date_interval)

    def mapper(self, line):
        for word in line.strip().split():
            yield word, 1

    def reducer(self, key, values):
        yield key, sum(values)
```

LUIGITARGETS

- HDFS
- Local File
- Postgres / MySQL, Redshift, ElasticSearch
- ... Easy to extend

DEFINING A TARGET

- Implement:

```
def exists(self)
```

And optionally:

```
connect or open / close
```

EXAMPLE

MYSQL TARGET

```
class MySqlTarget(luigi.Target):

    def touch(self, connection=None):
        ...

    def exists(self, connection=None):
        cursor = connection.cursor()
        cursor.execute("""SELECT 1 FROM {marker_table}
                           WHERE update_id = %s
                           LIMIT 1""".format(marker_table=self.marker_table),
                      (self.update_id,))
        row = cursor.fetchone()
        return row is not None

    def connect(self, autocommit=False):
        ...

    def create_marker_table(self):
        ...
```

THE GRAND SCHEME

THE GRAND SCHEME

**Run
Task**

THE GRAND SCHEME

**Run
Task**

**Check
Deps**

THE GRAND SCHEME

**Run
Task**

`self.requires()`

**Check
Deps**

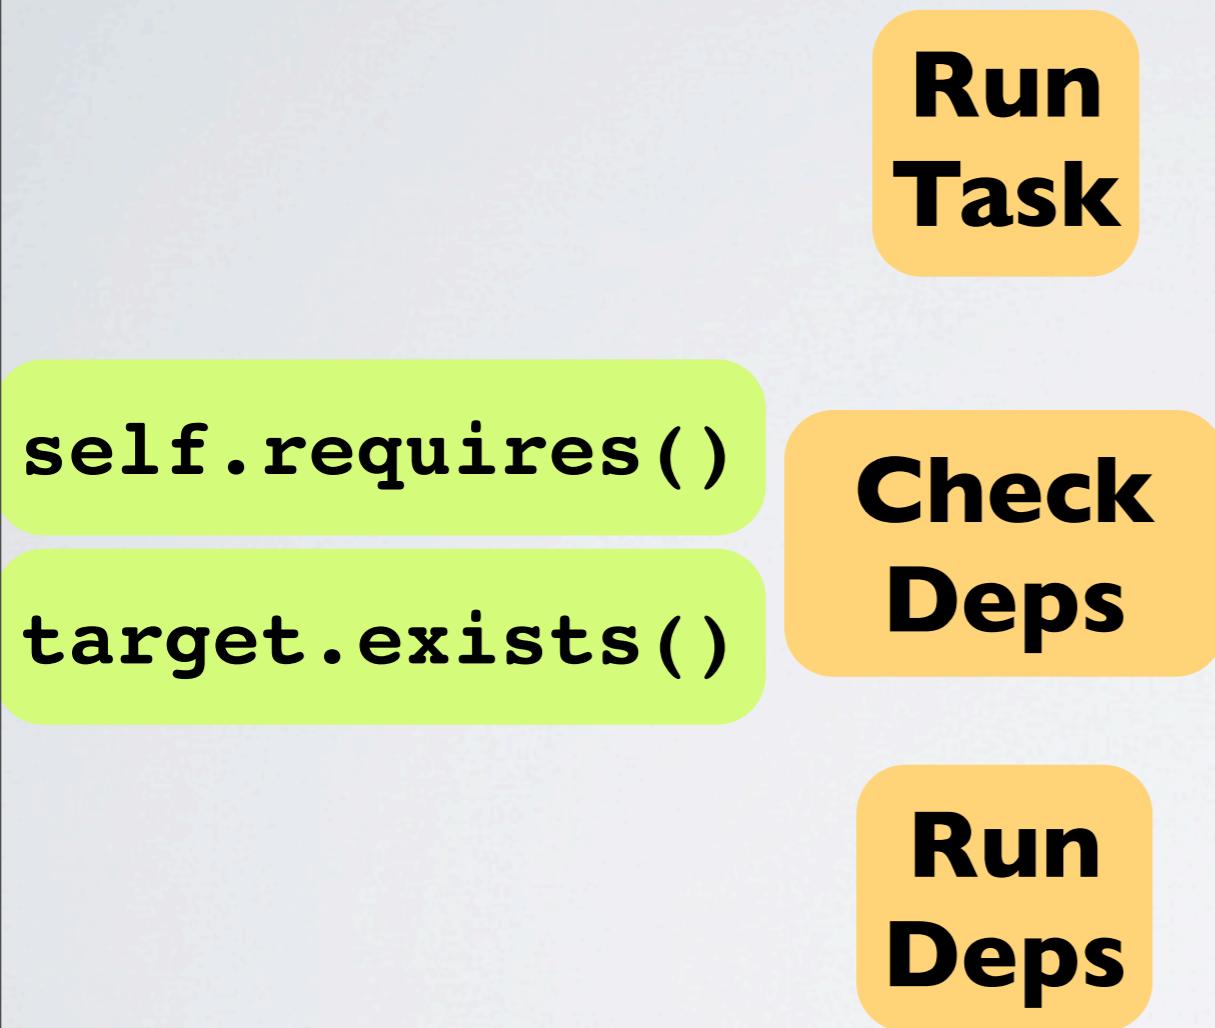
THE GRAND SCHEME

**Run
Task**

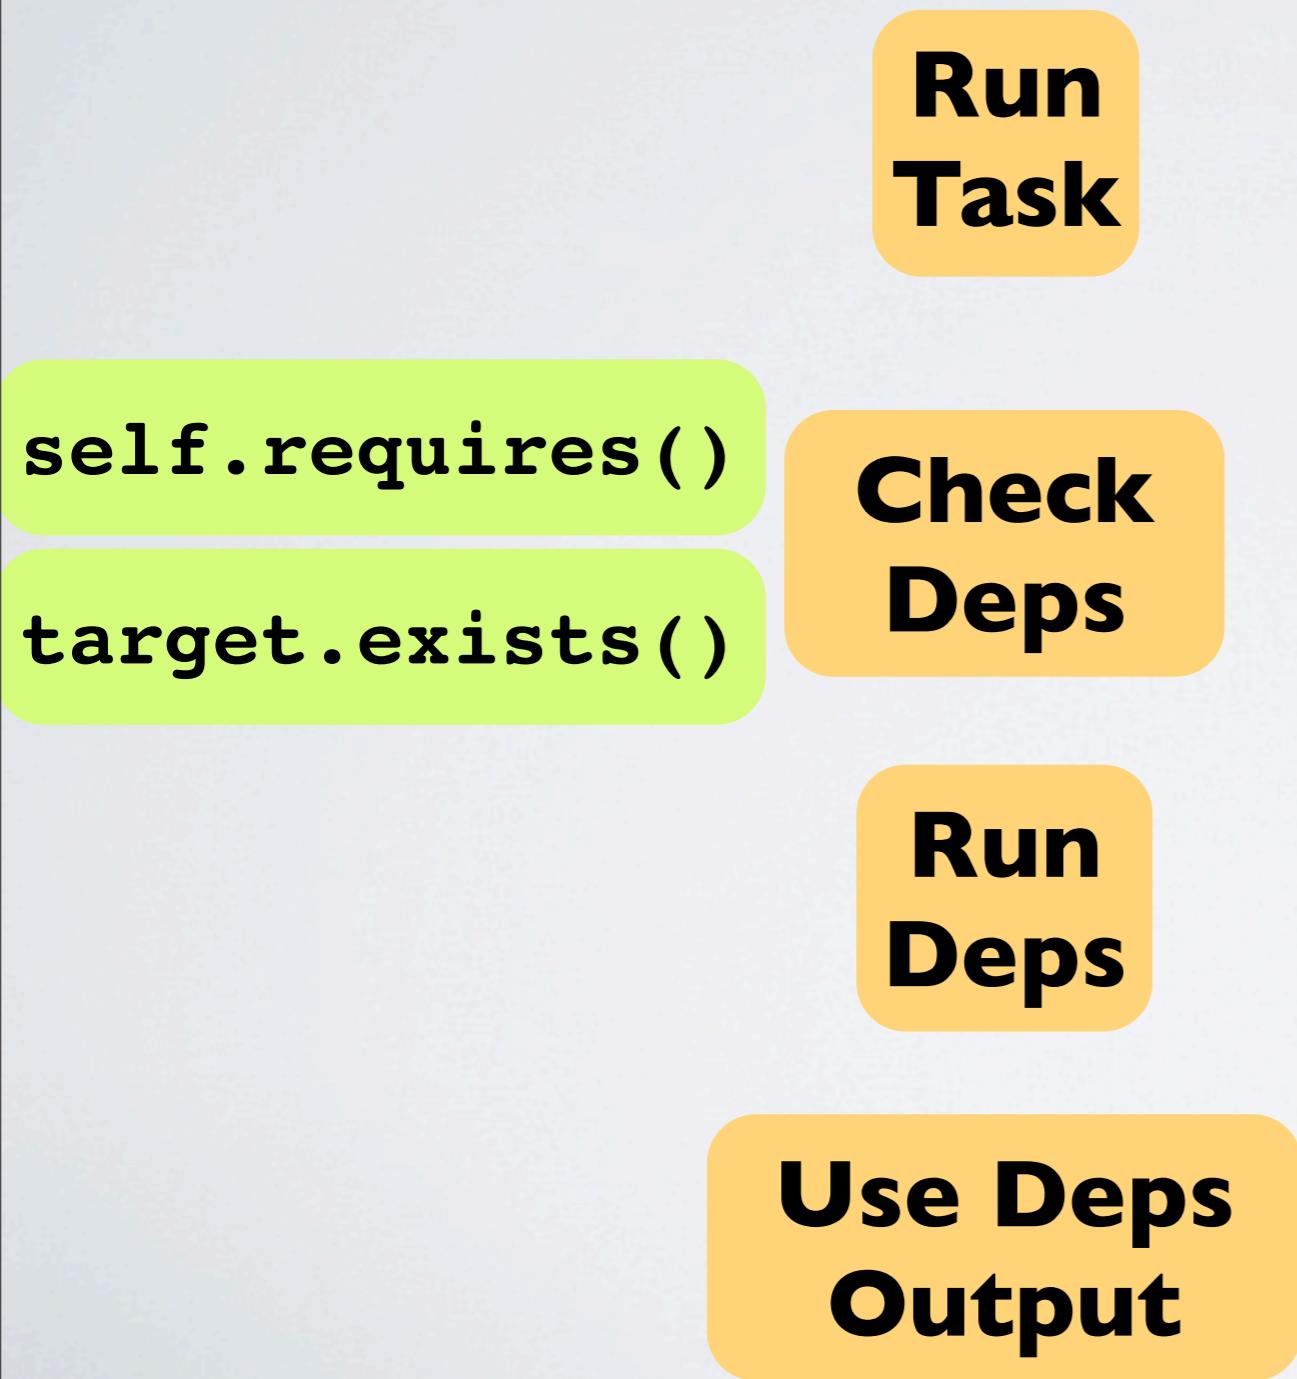
```
self.requires()  
target.exists()
```

**Check
Deps**

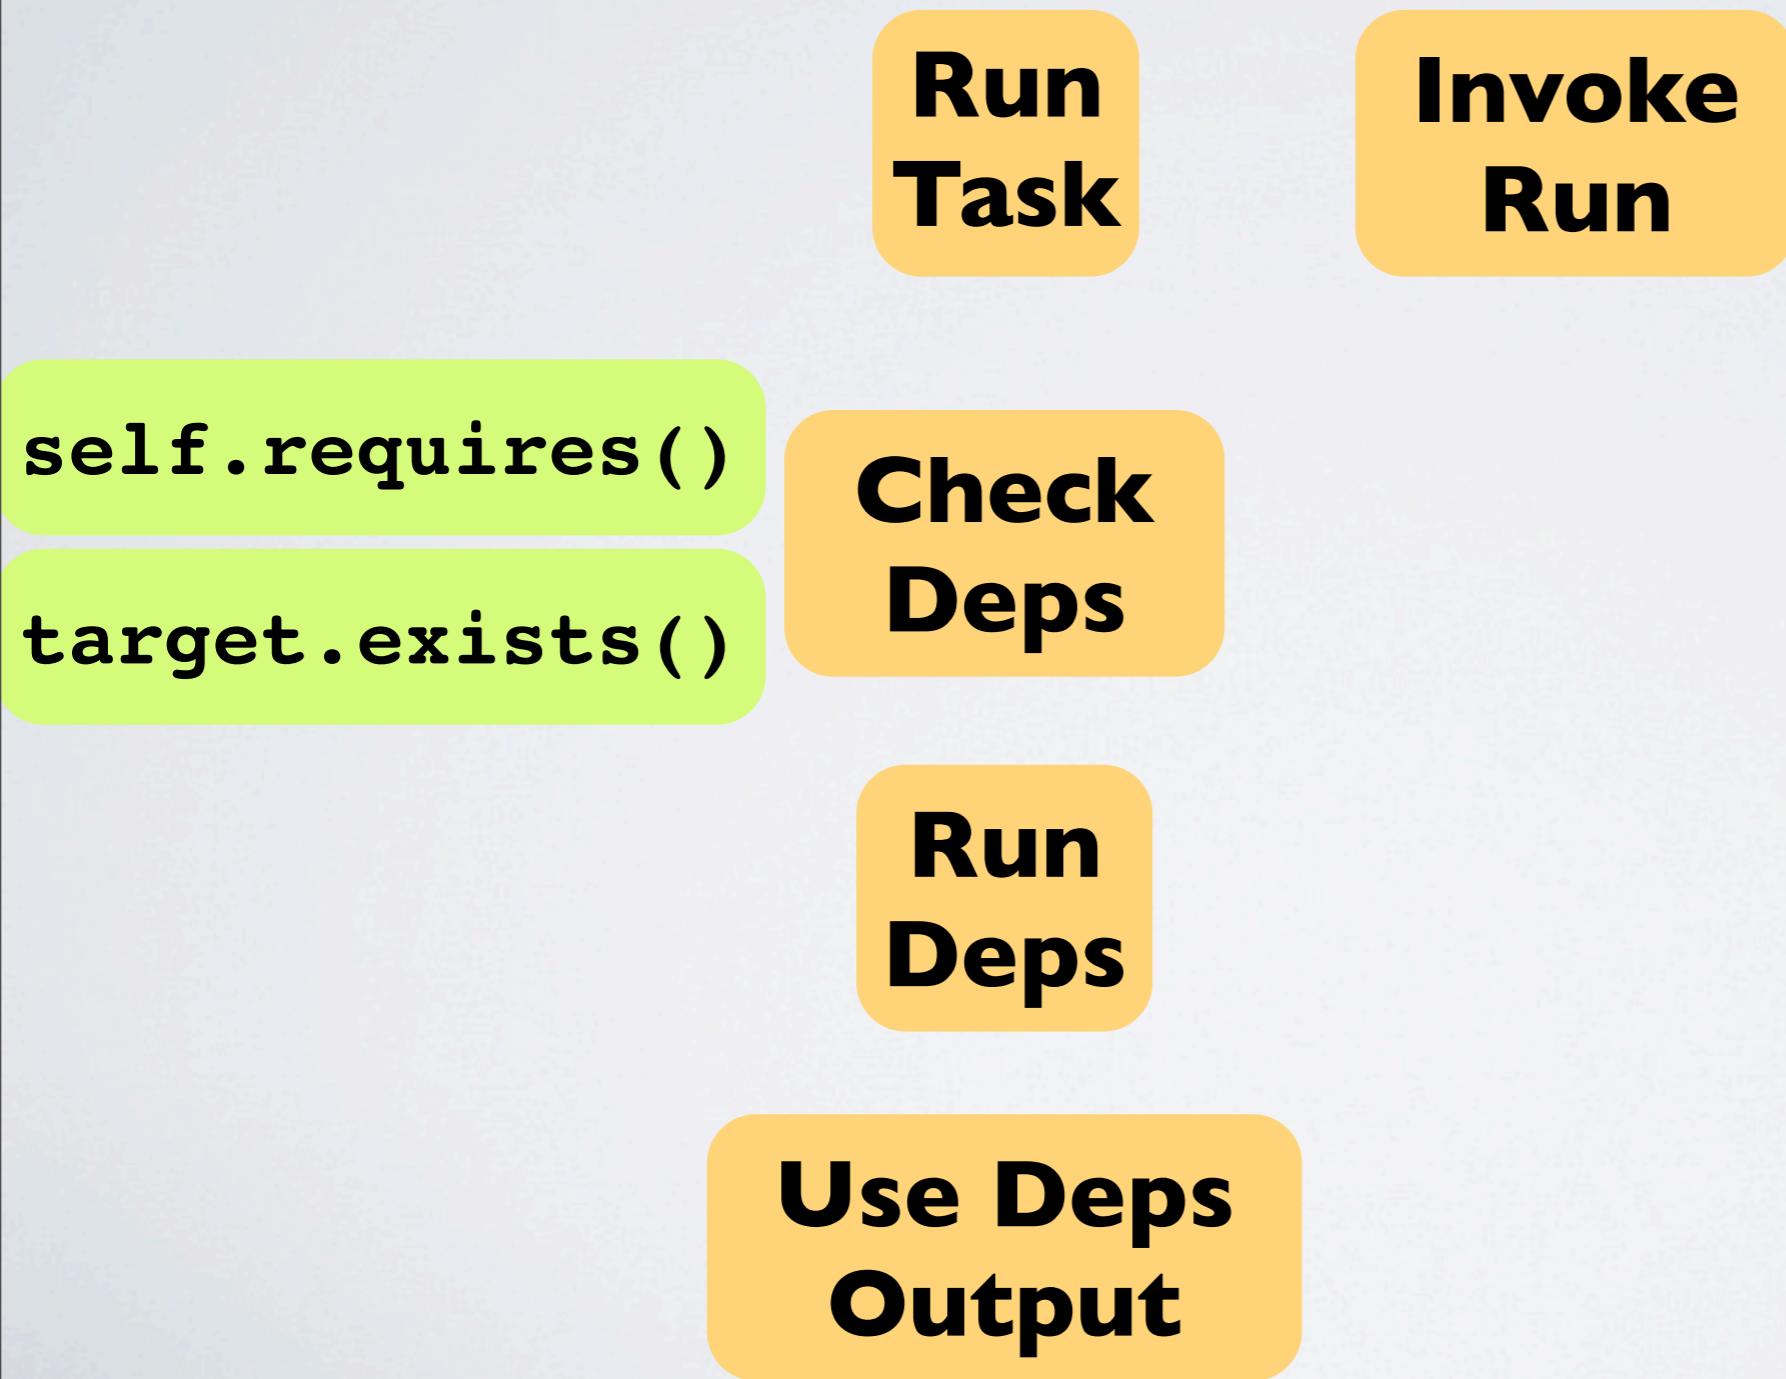
THE GRAND SCHEME



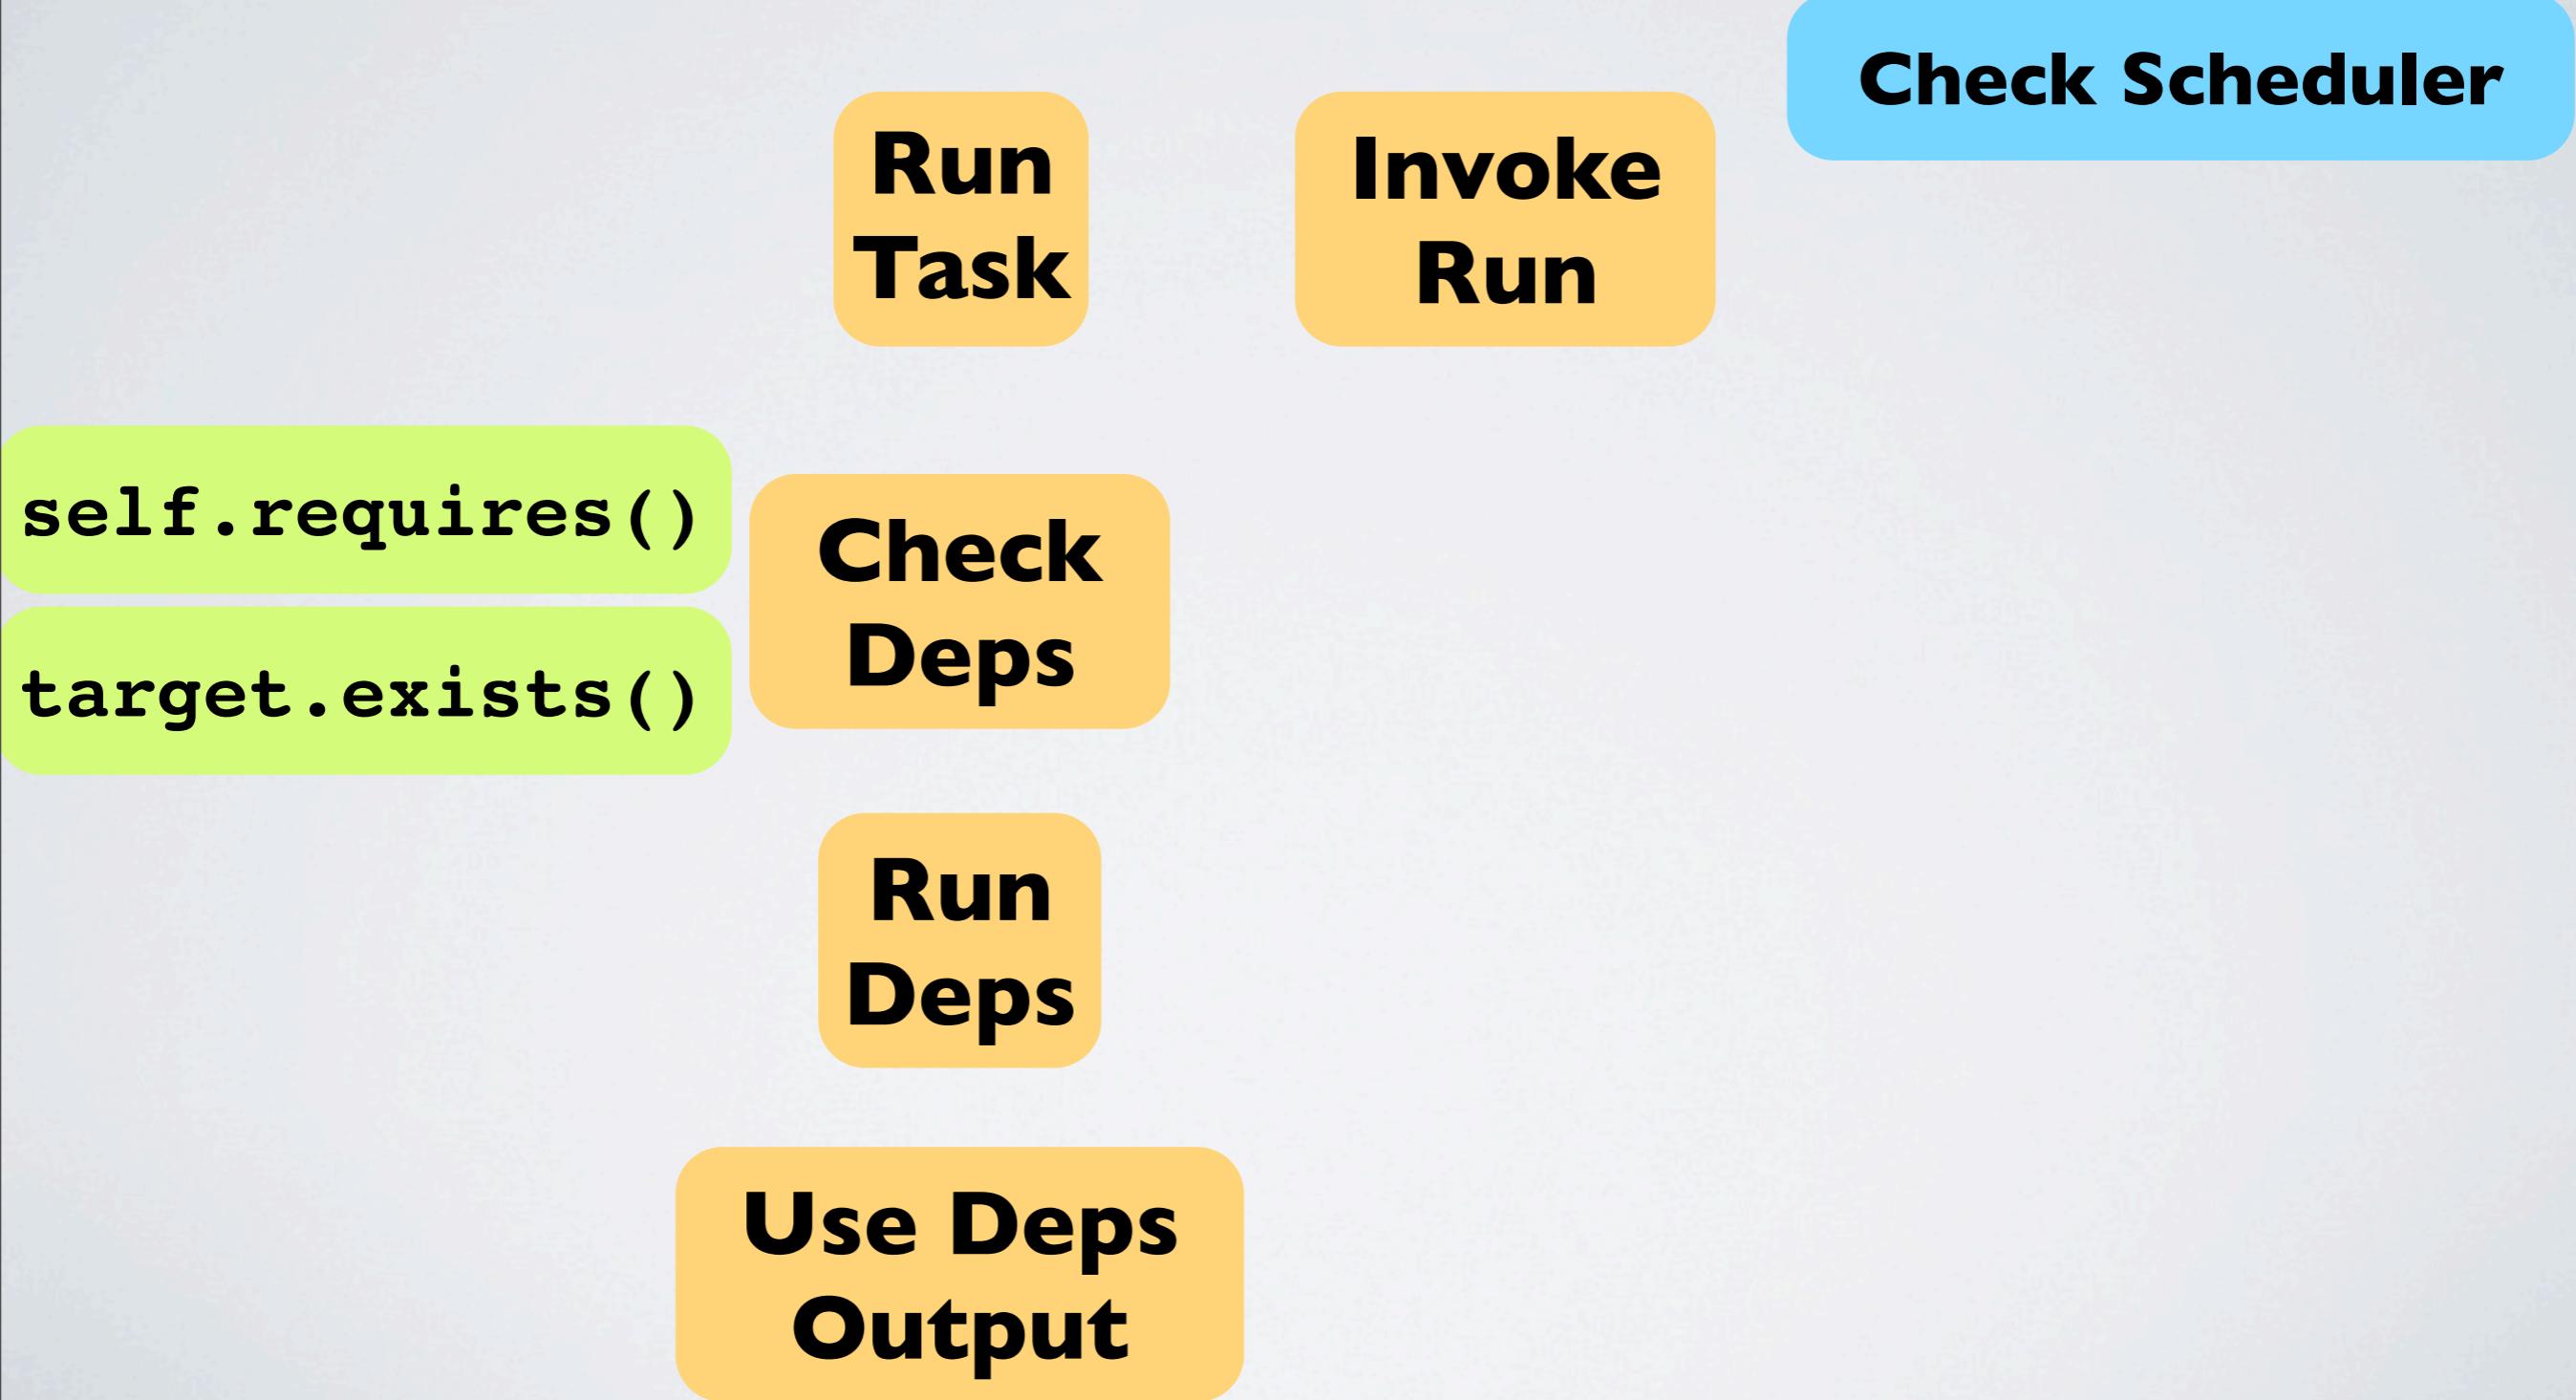
THE GRAND SCHEME



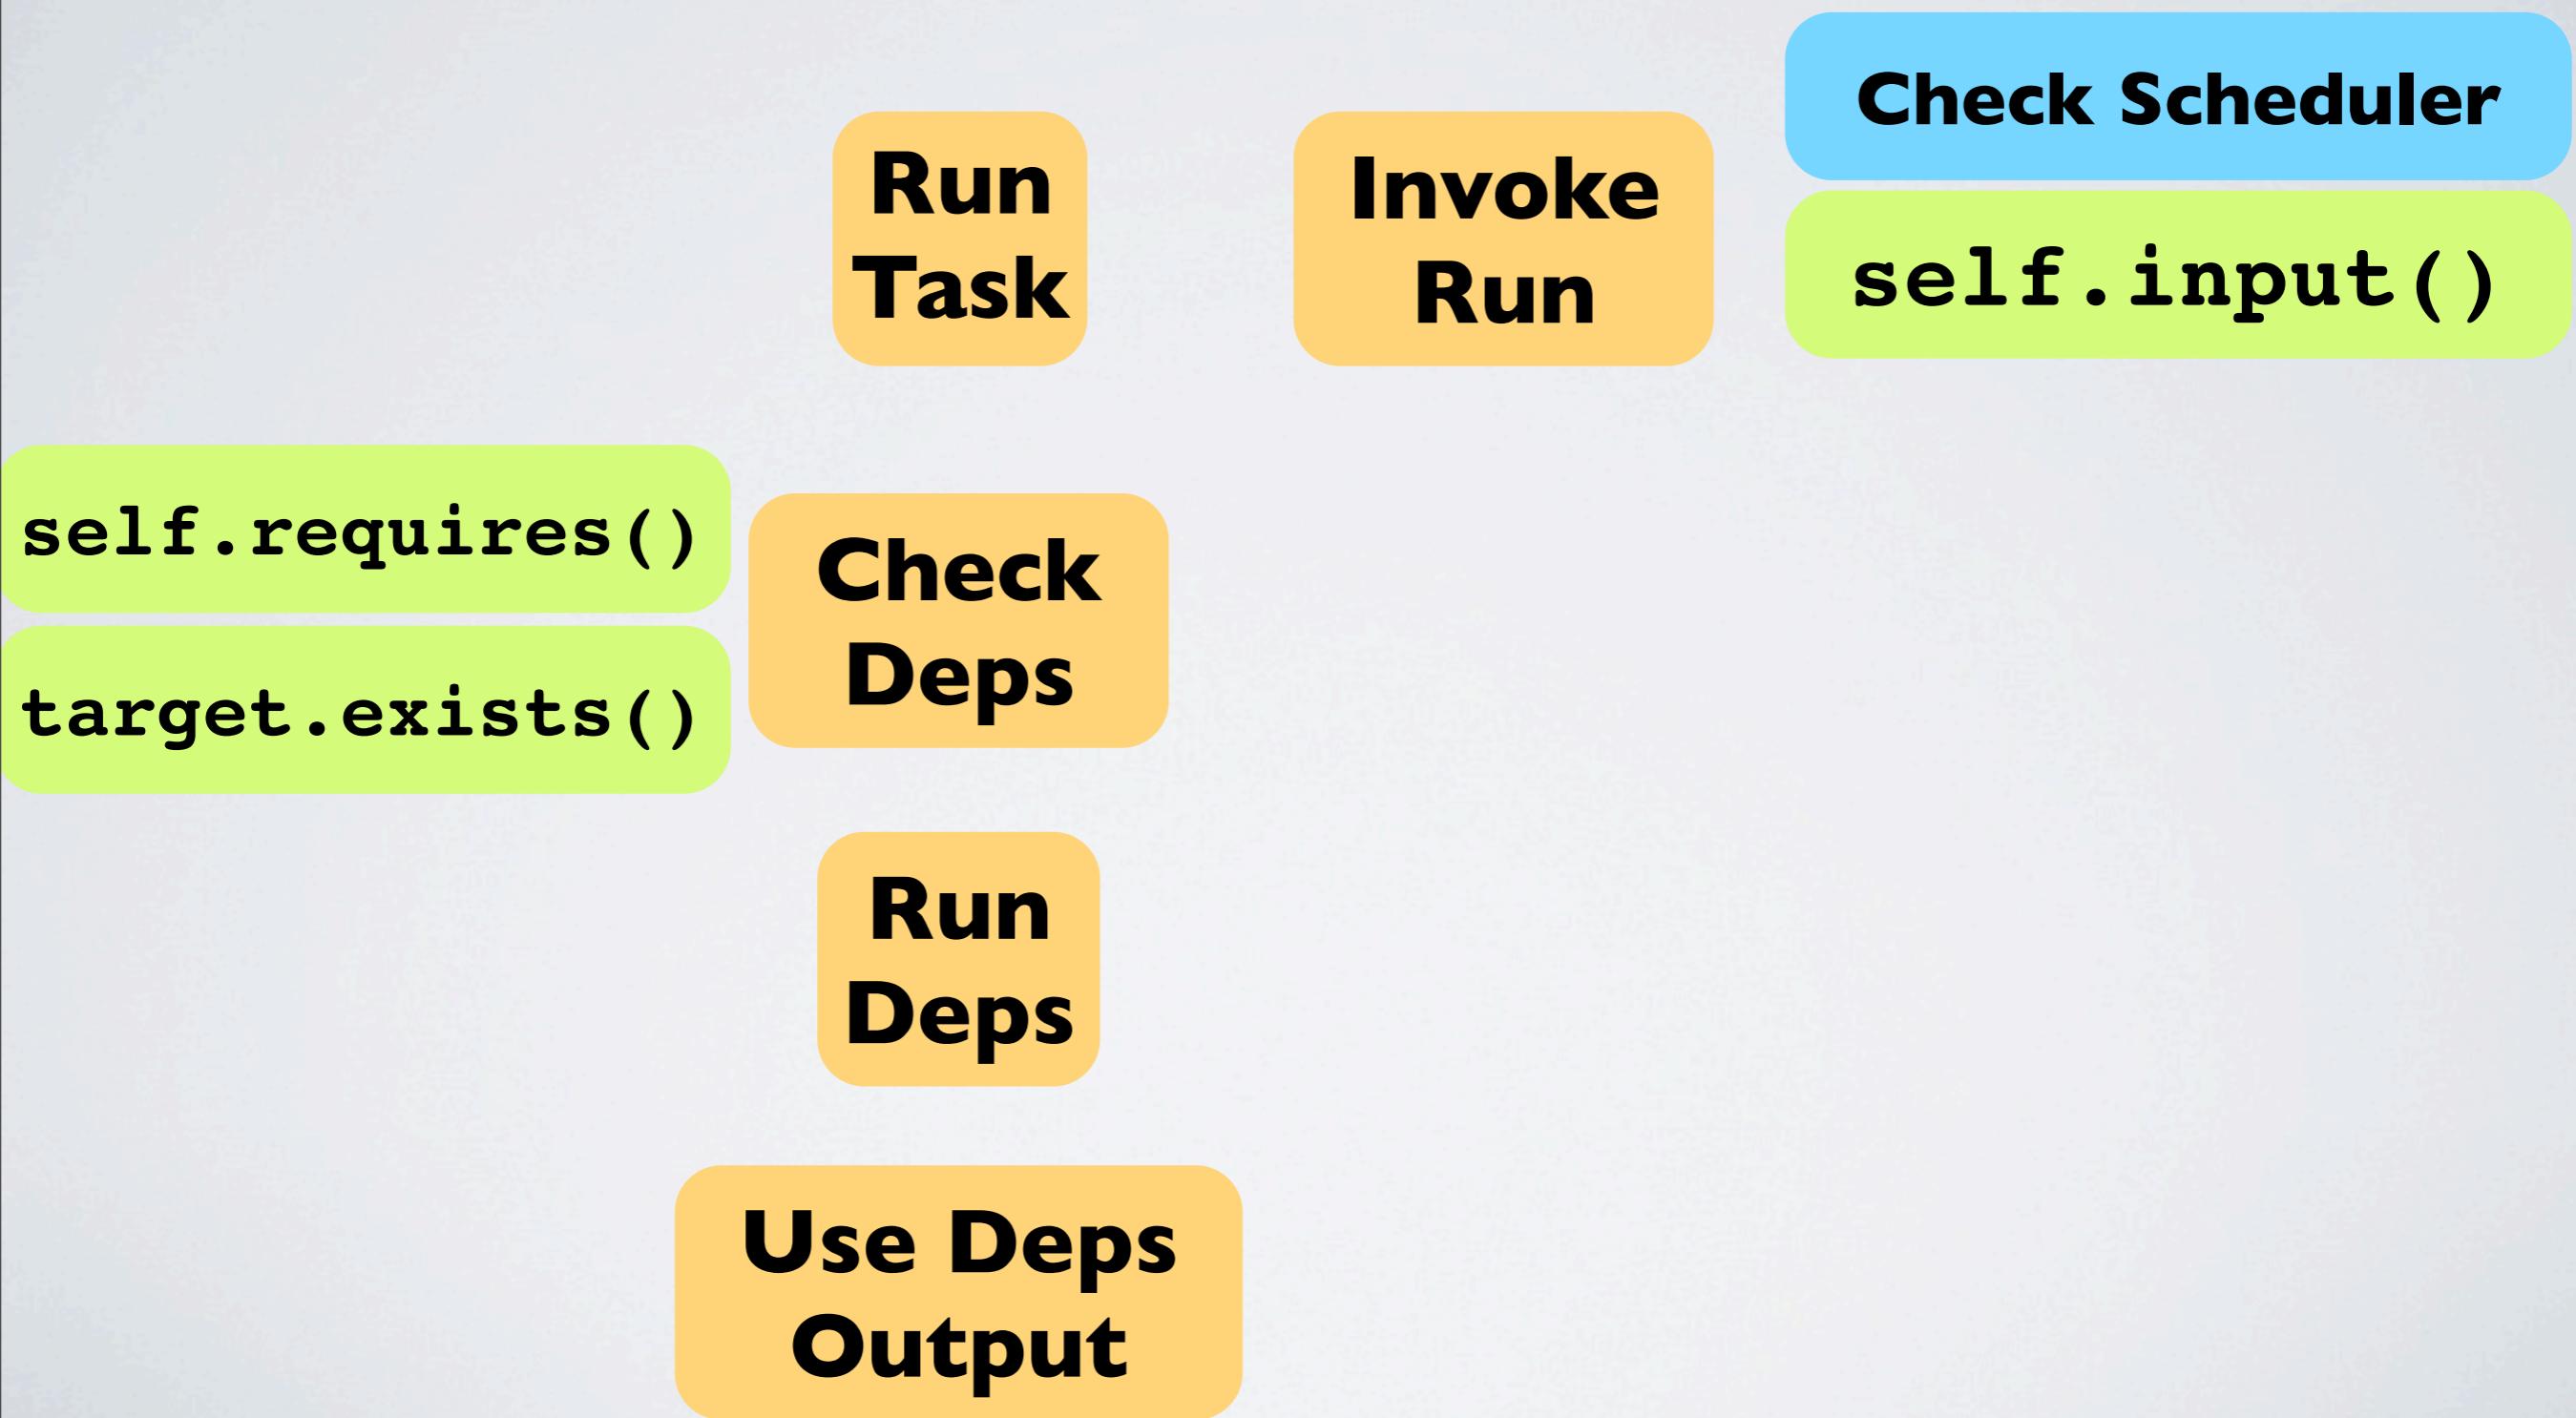
THE GRAND SCHEME



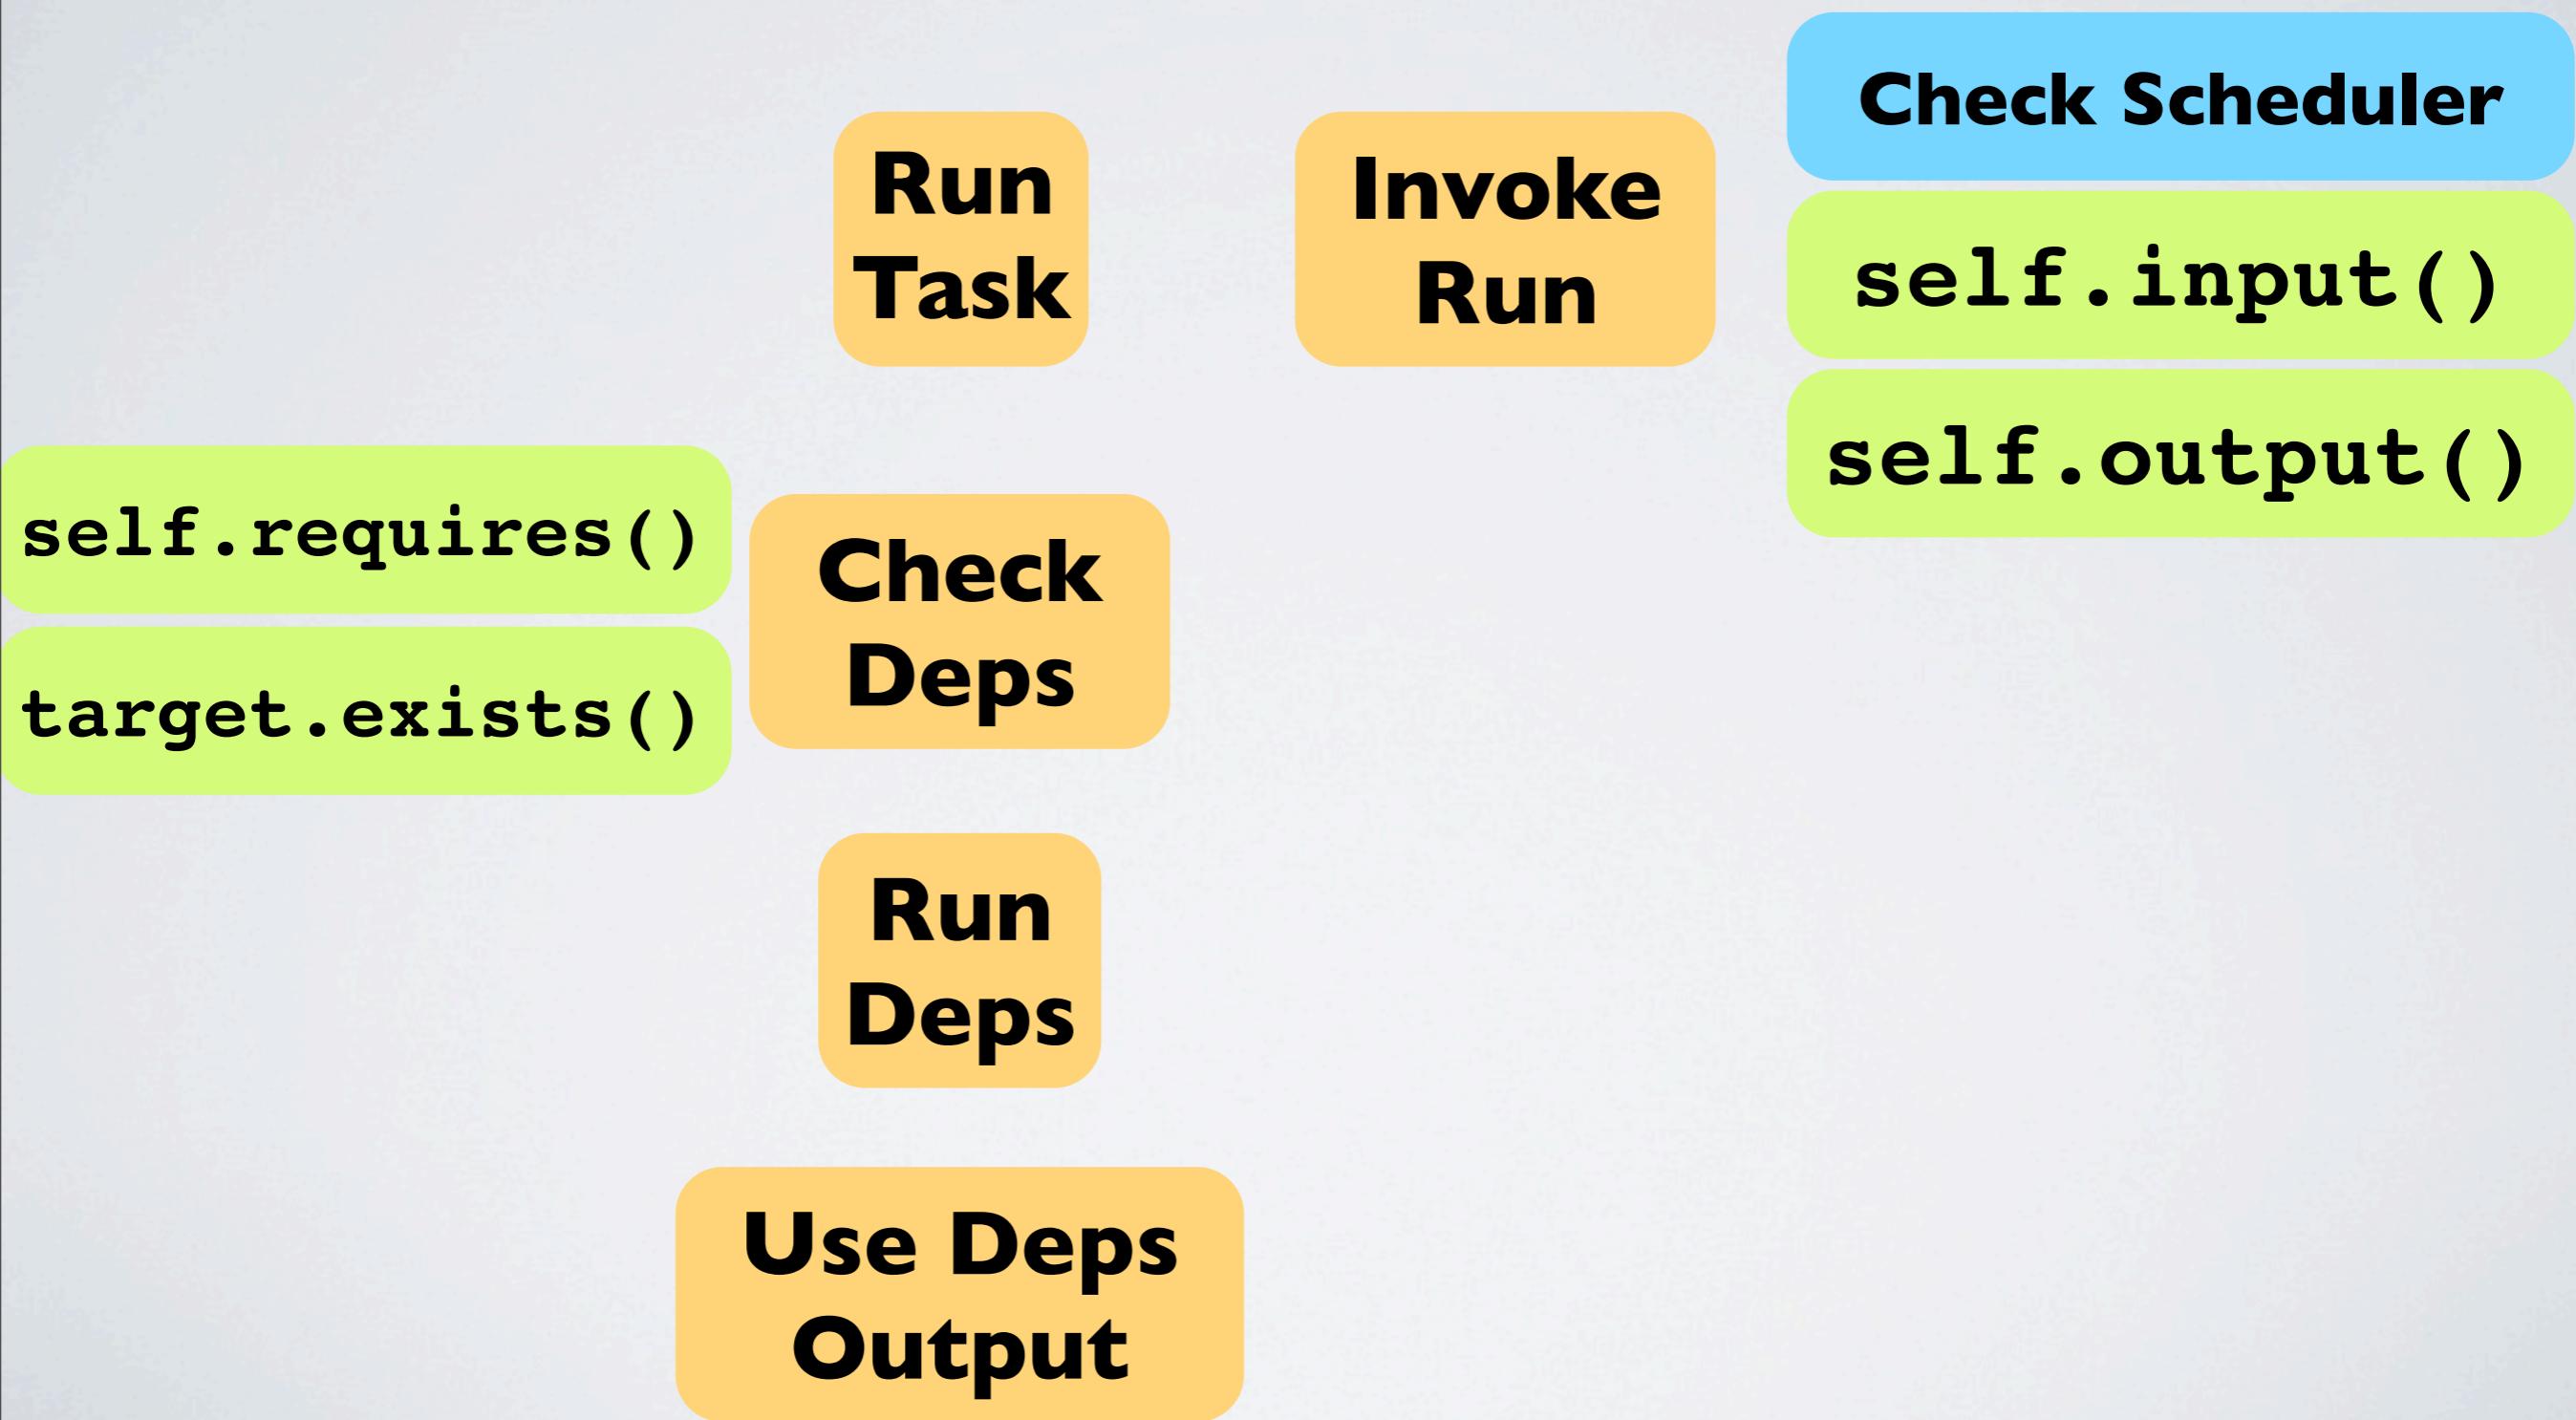
THE GRAND SCHEME



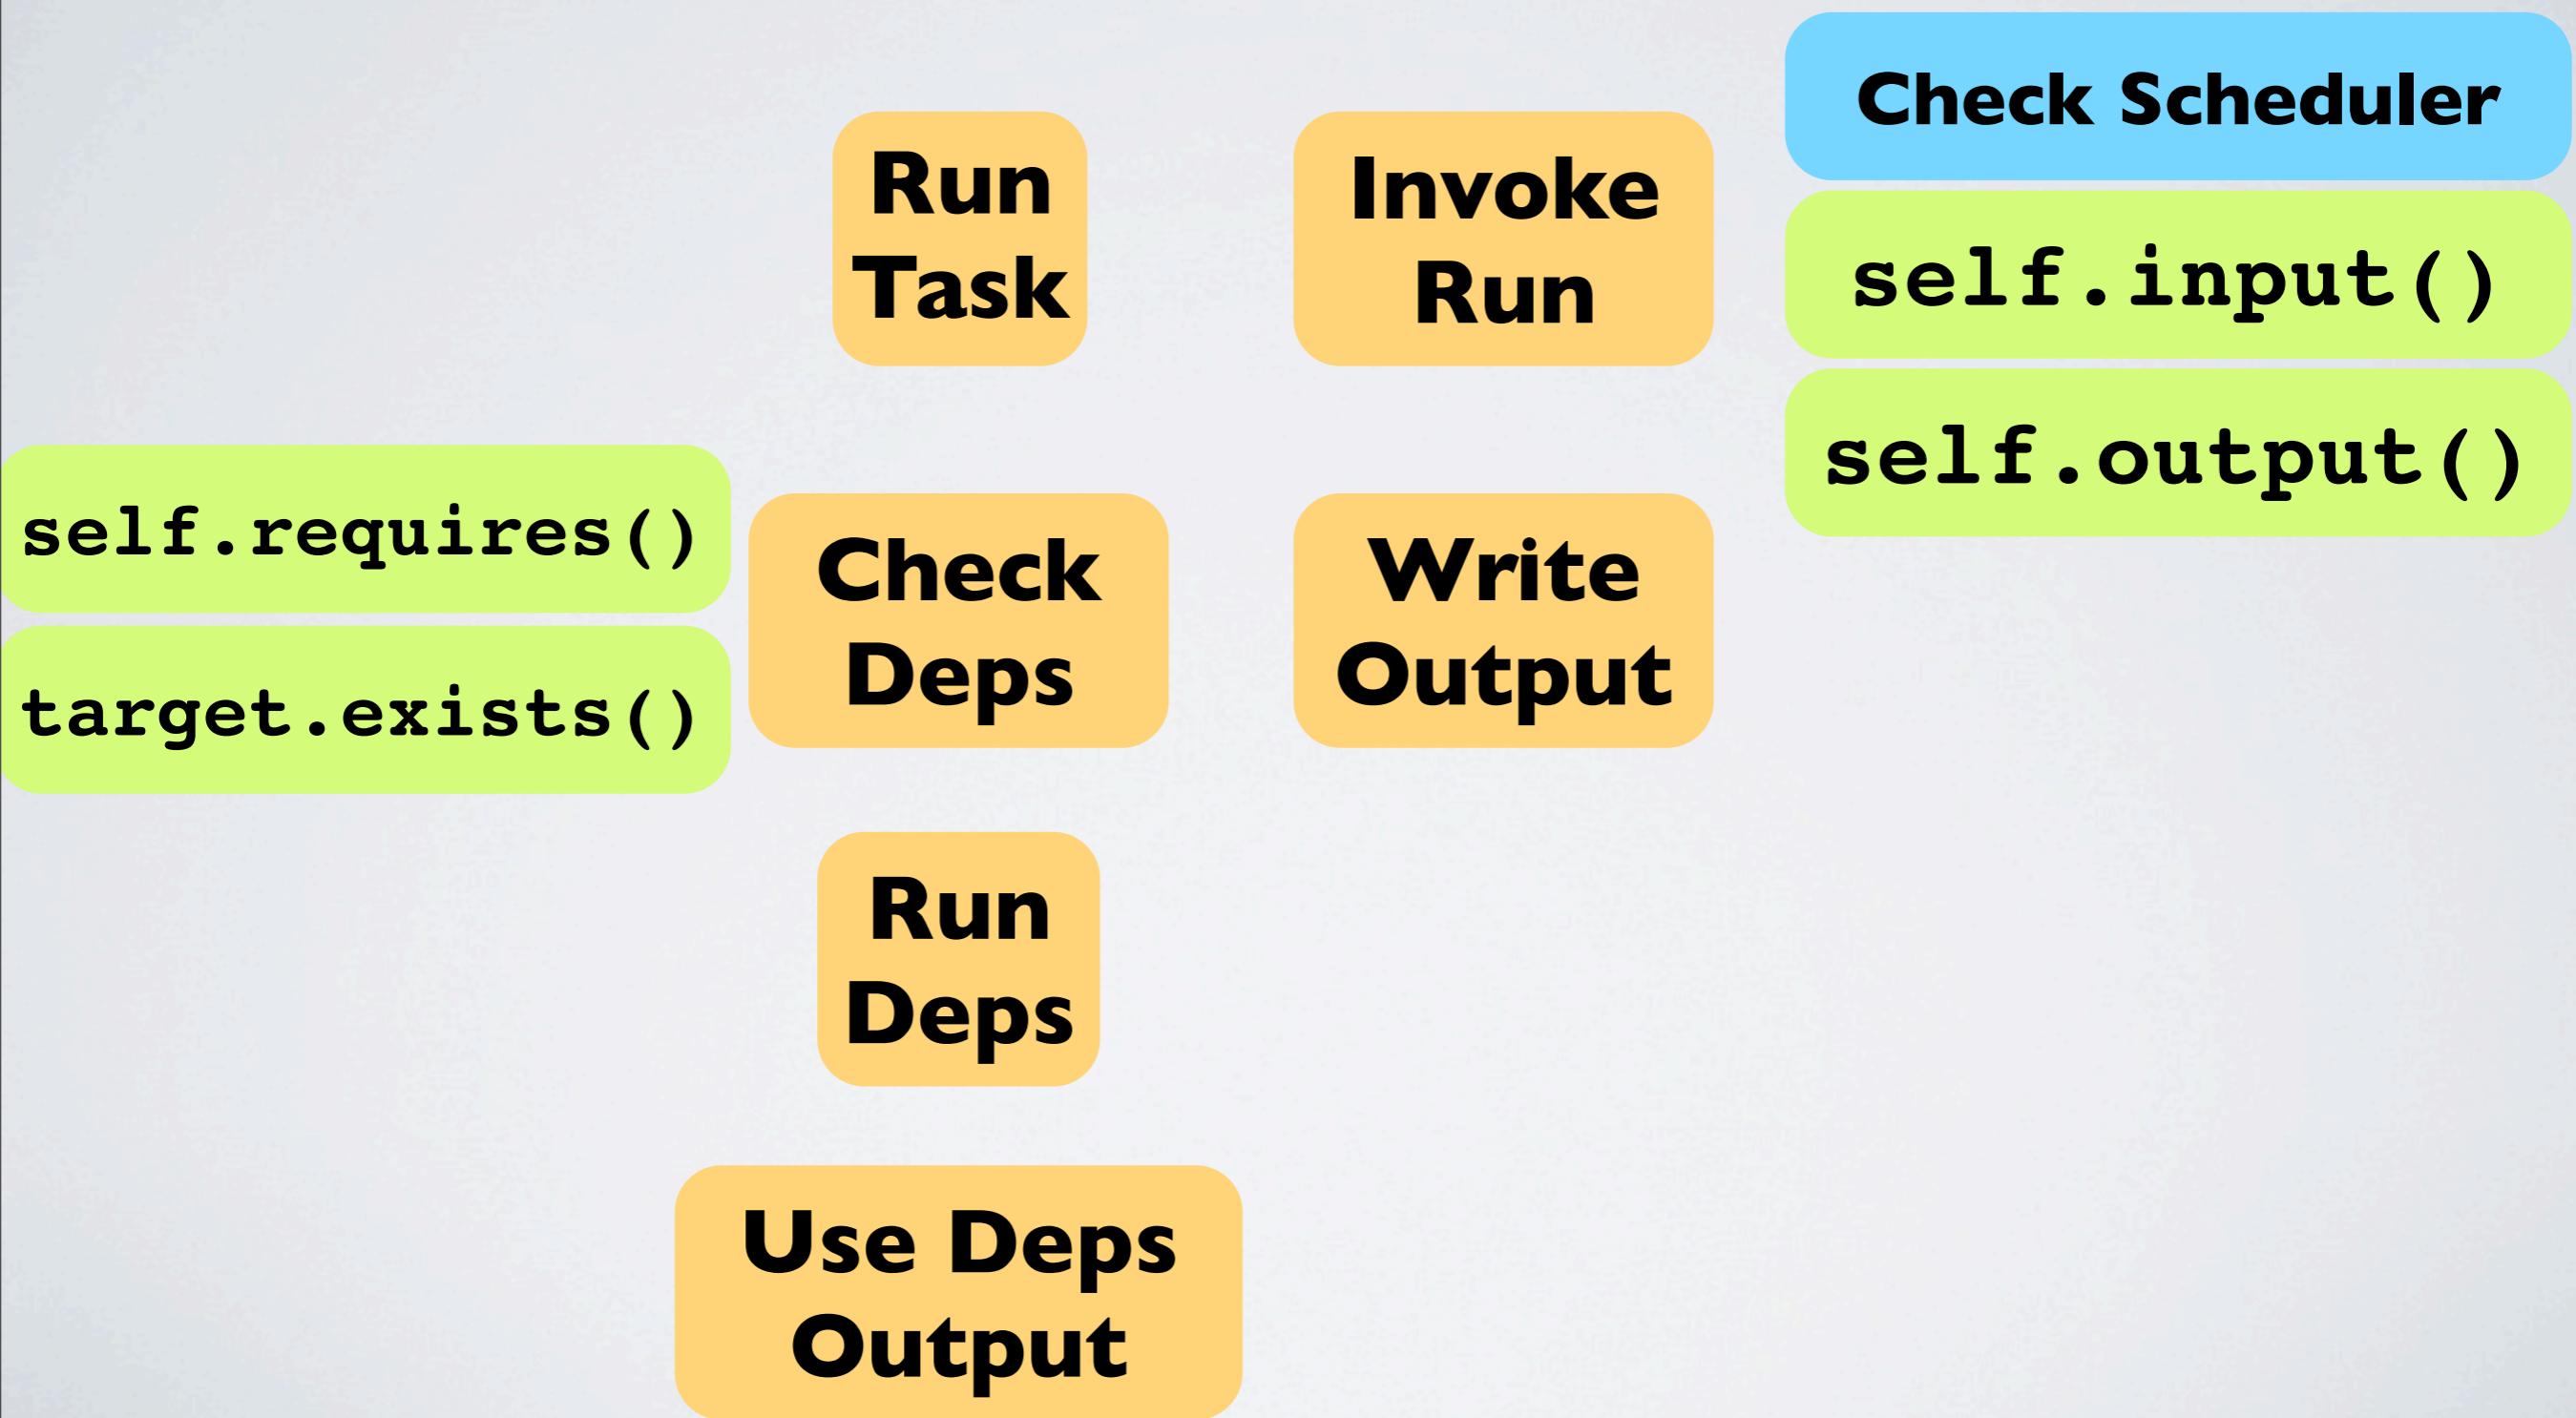
THE GRAND SCHEME



THE GRAND SCHEME



THE GRAND SCHEME

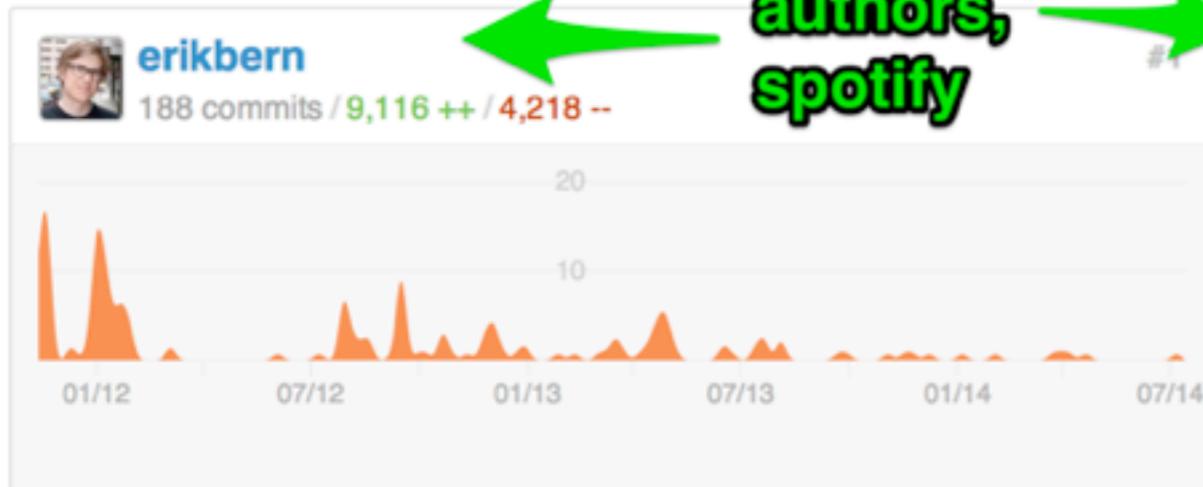
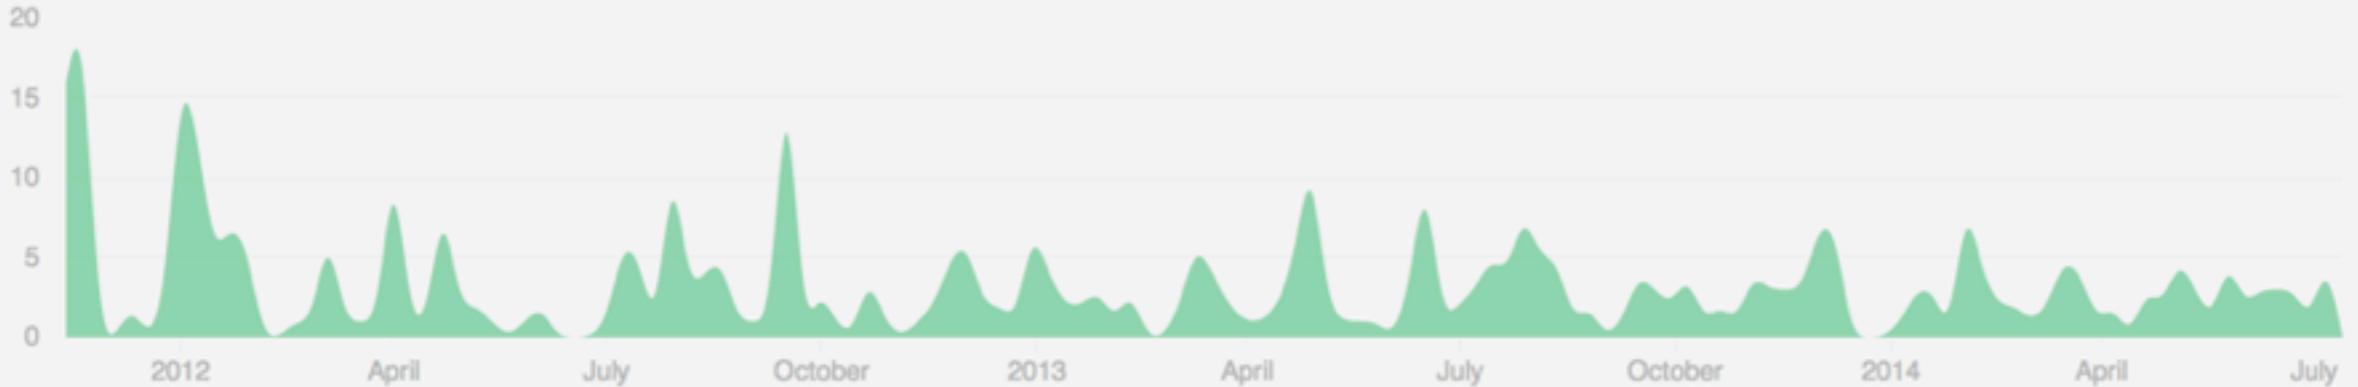


OPEN SOURCE

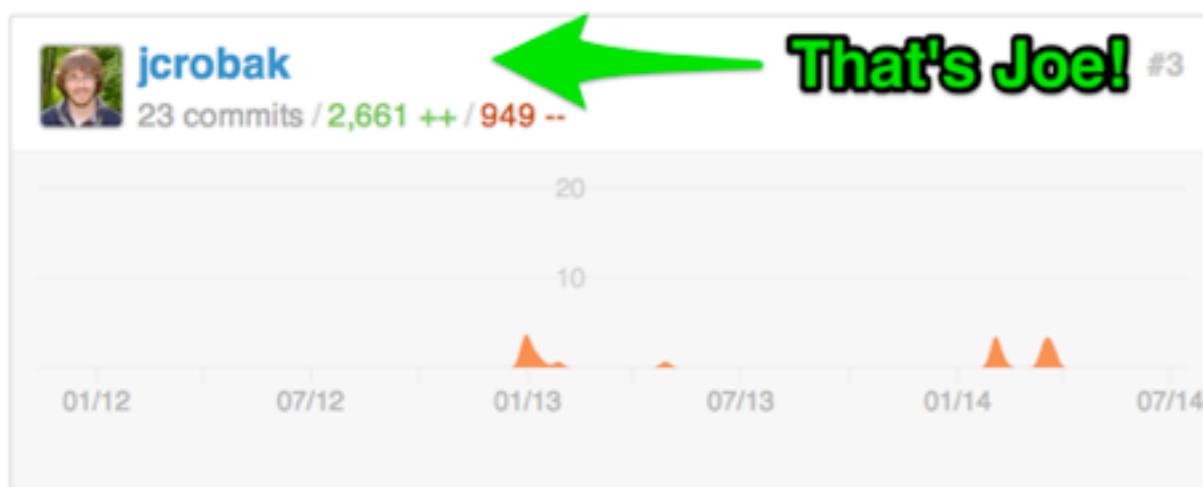
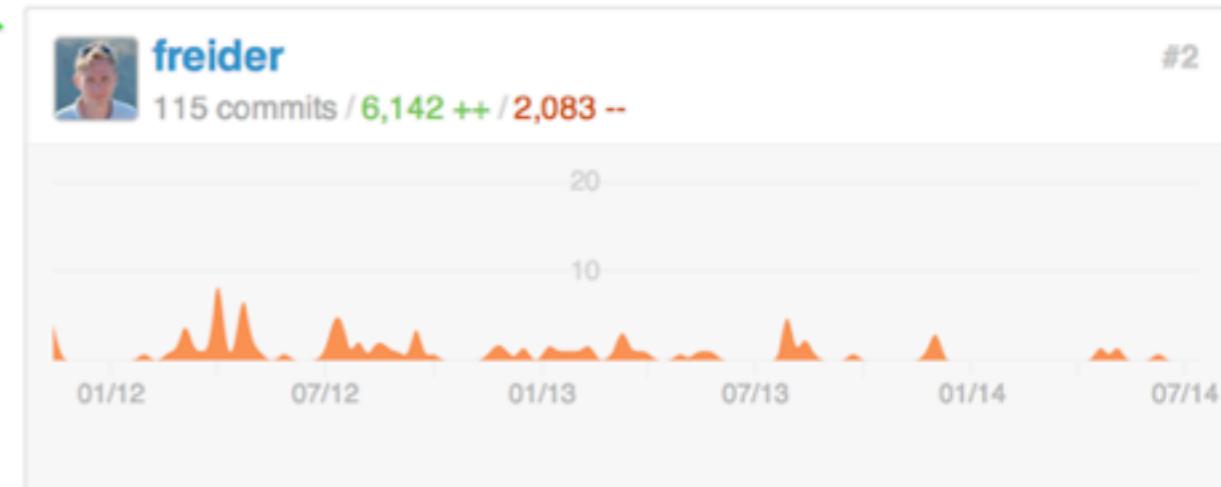


BY

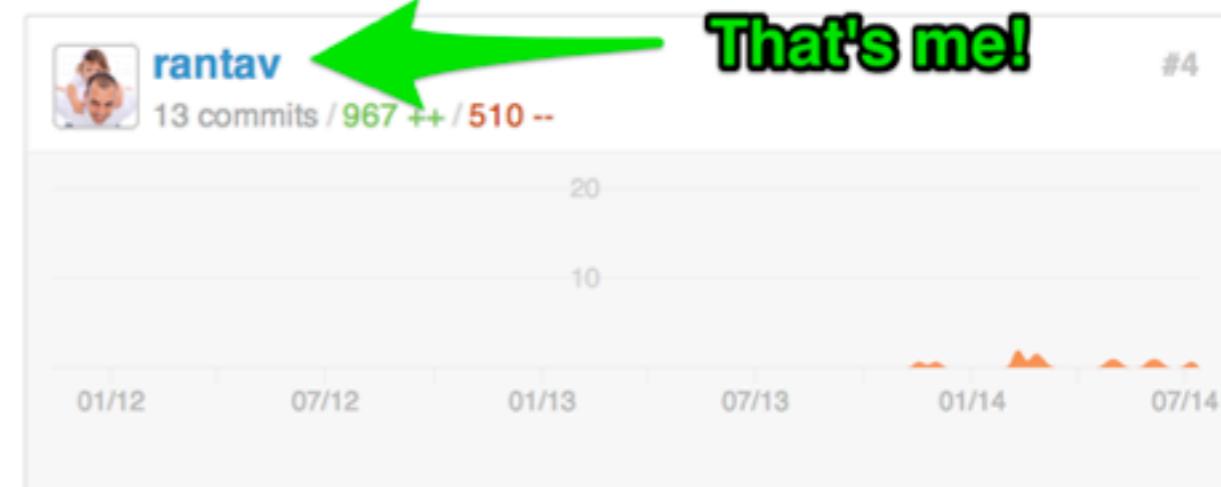




← authors,
spotify



← That's Joe!



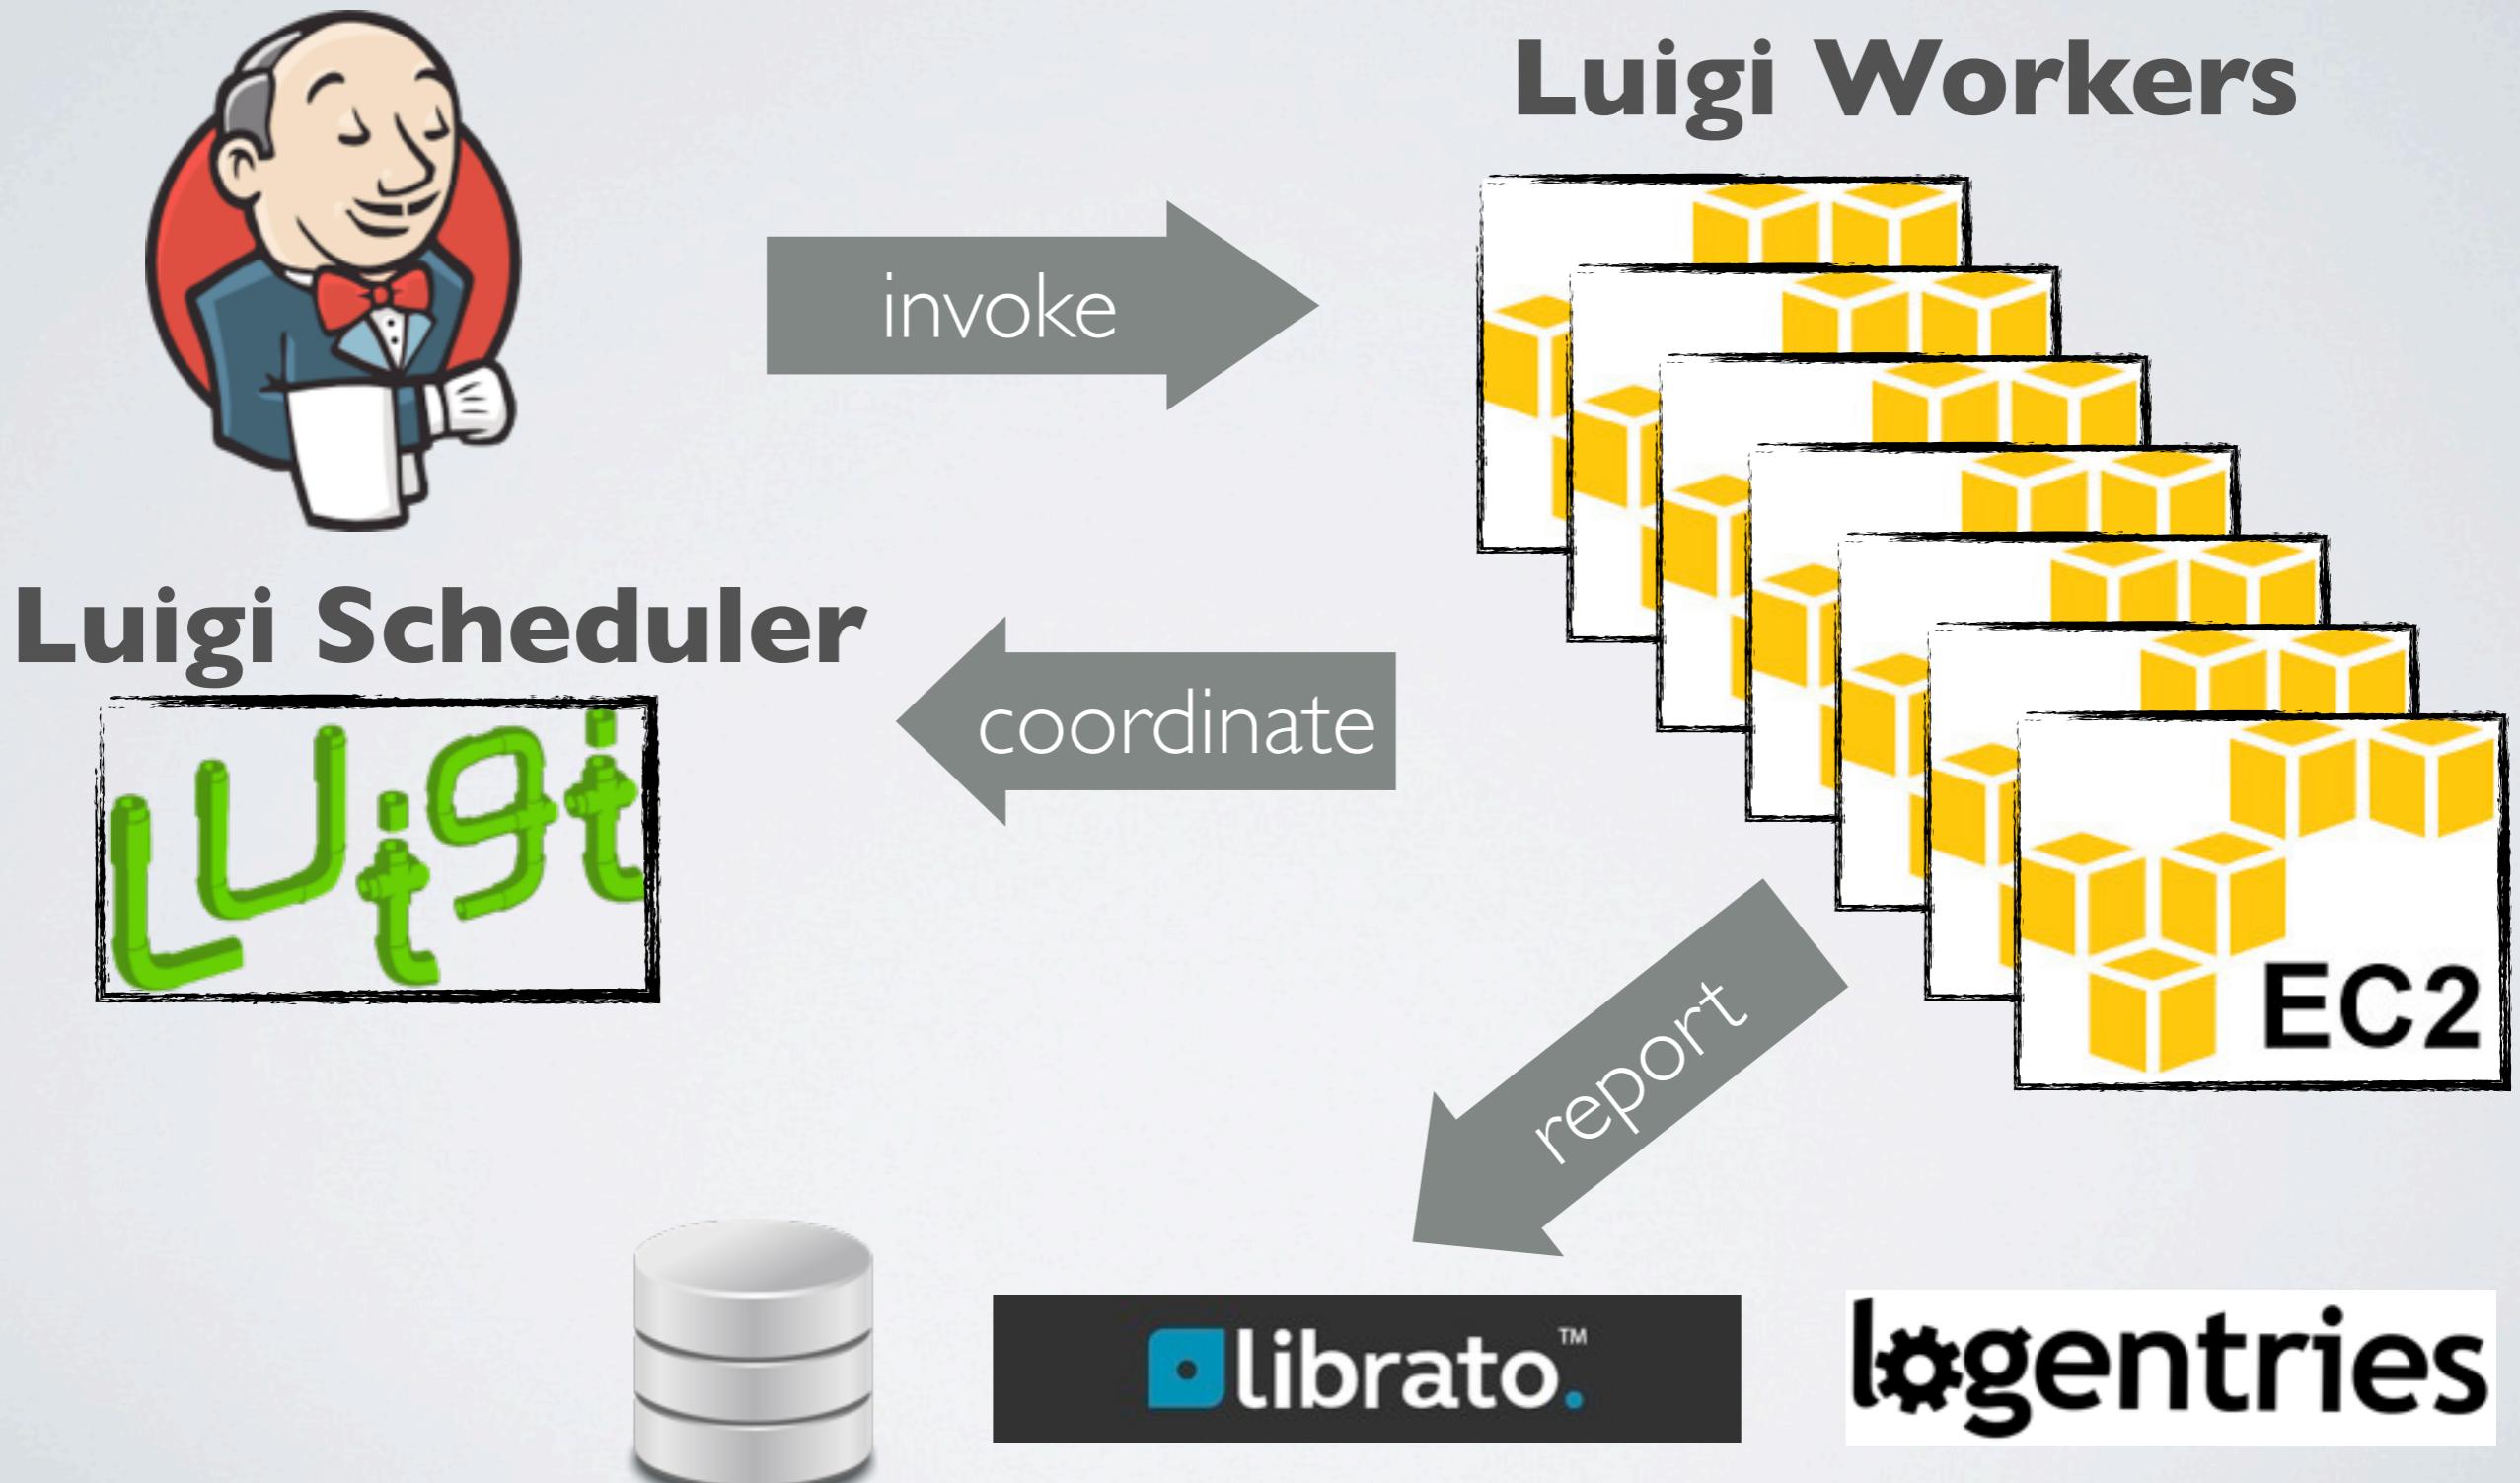
WHAT DID I DO?

- Add **Redshift** support
- Add **MySQL** support
- Various small features (improved **notifications**, **dep.py**, **historydb** etc)
- Various bug reports
 - And fixes!

LUIGI @ TOTANGO

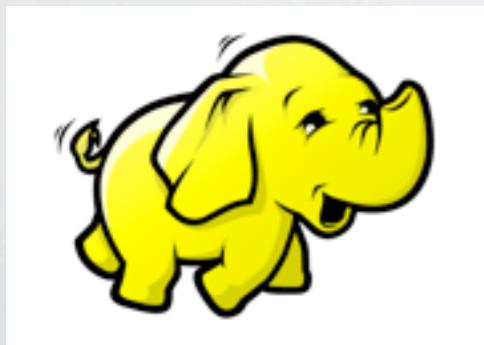
- Daily computation
- Hourly computation
- Ad-hoc data loading (for data analysis activities, to redshift)

TOTANGO'S SETUP

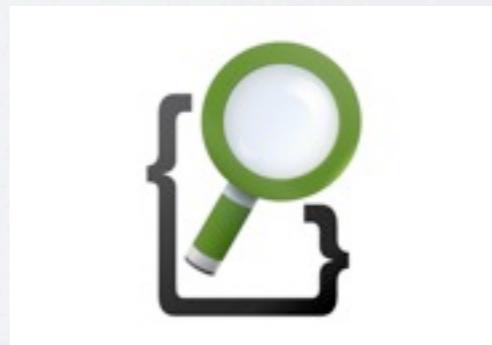


TOTANGO'S SETUP

Luigi Worker



Spark

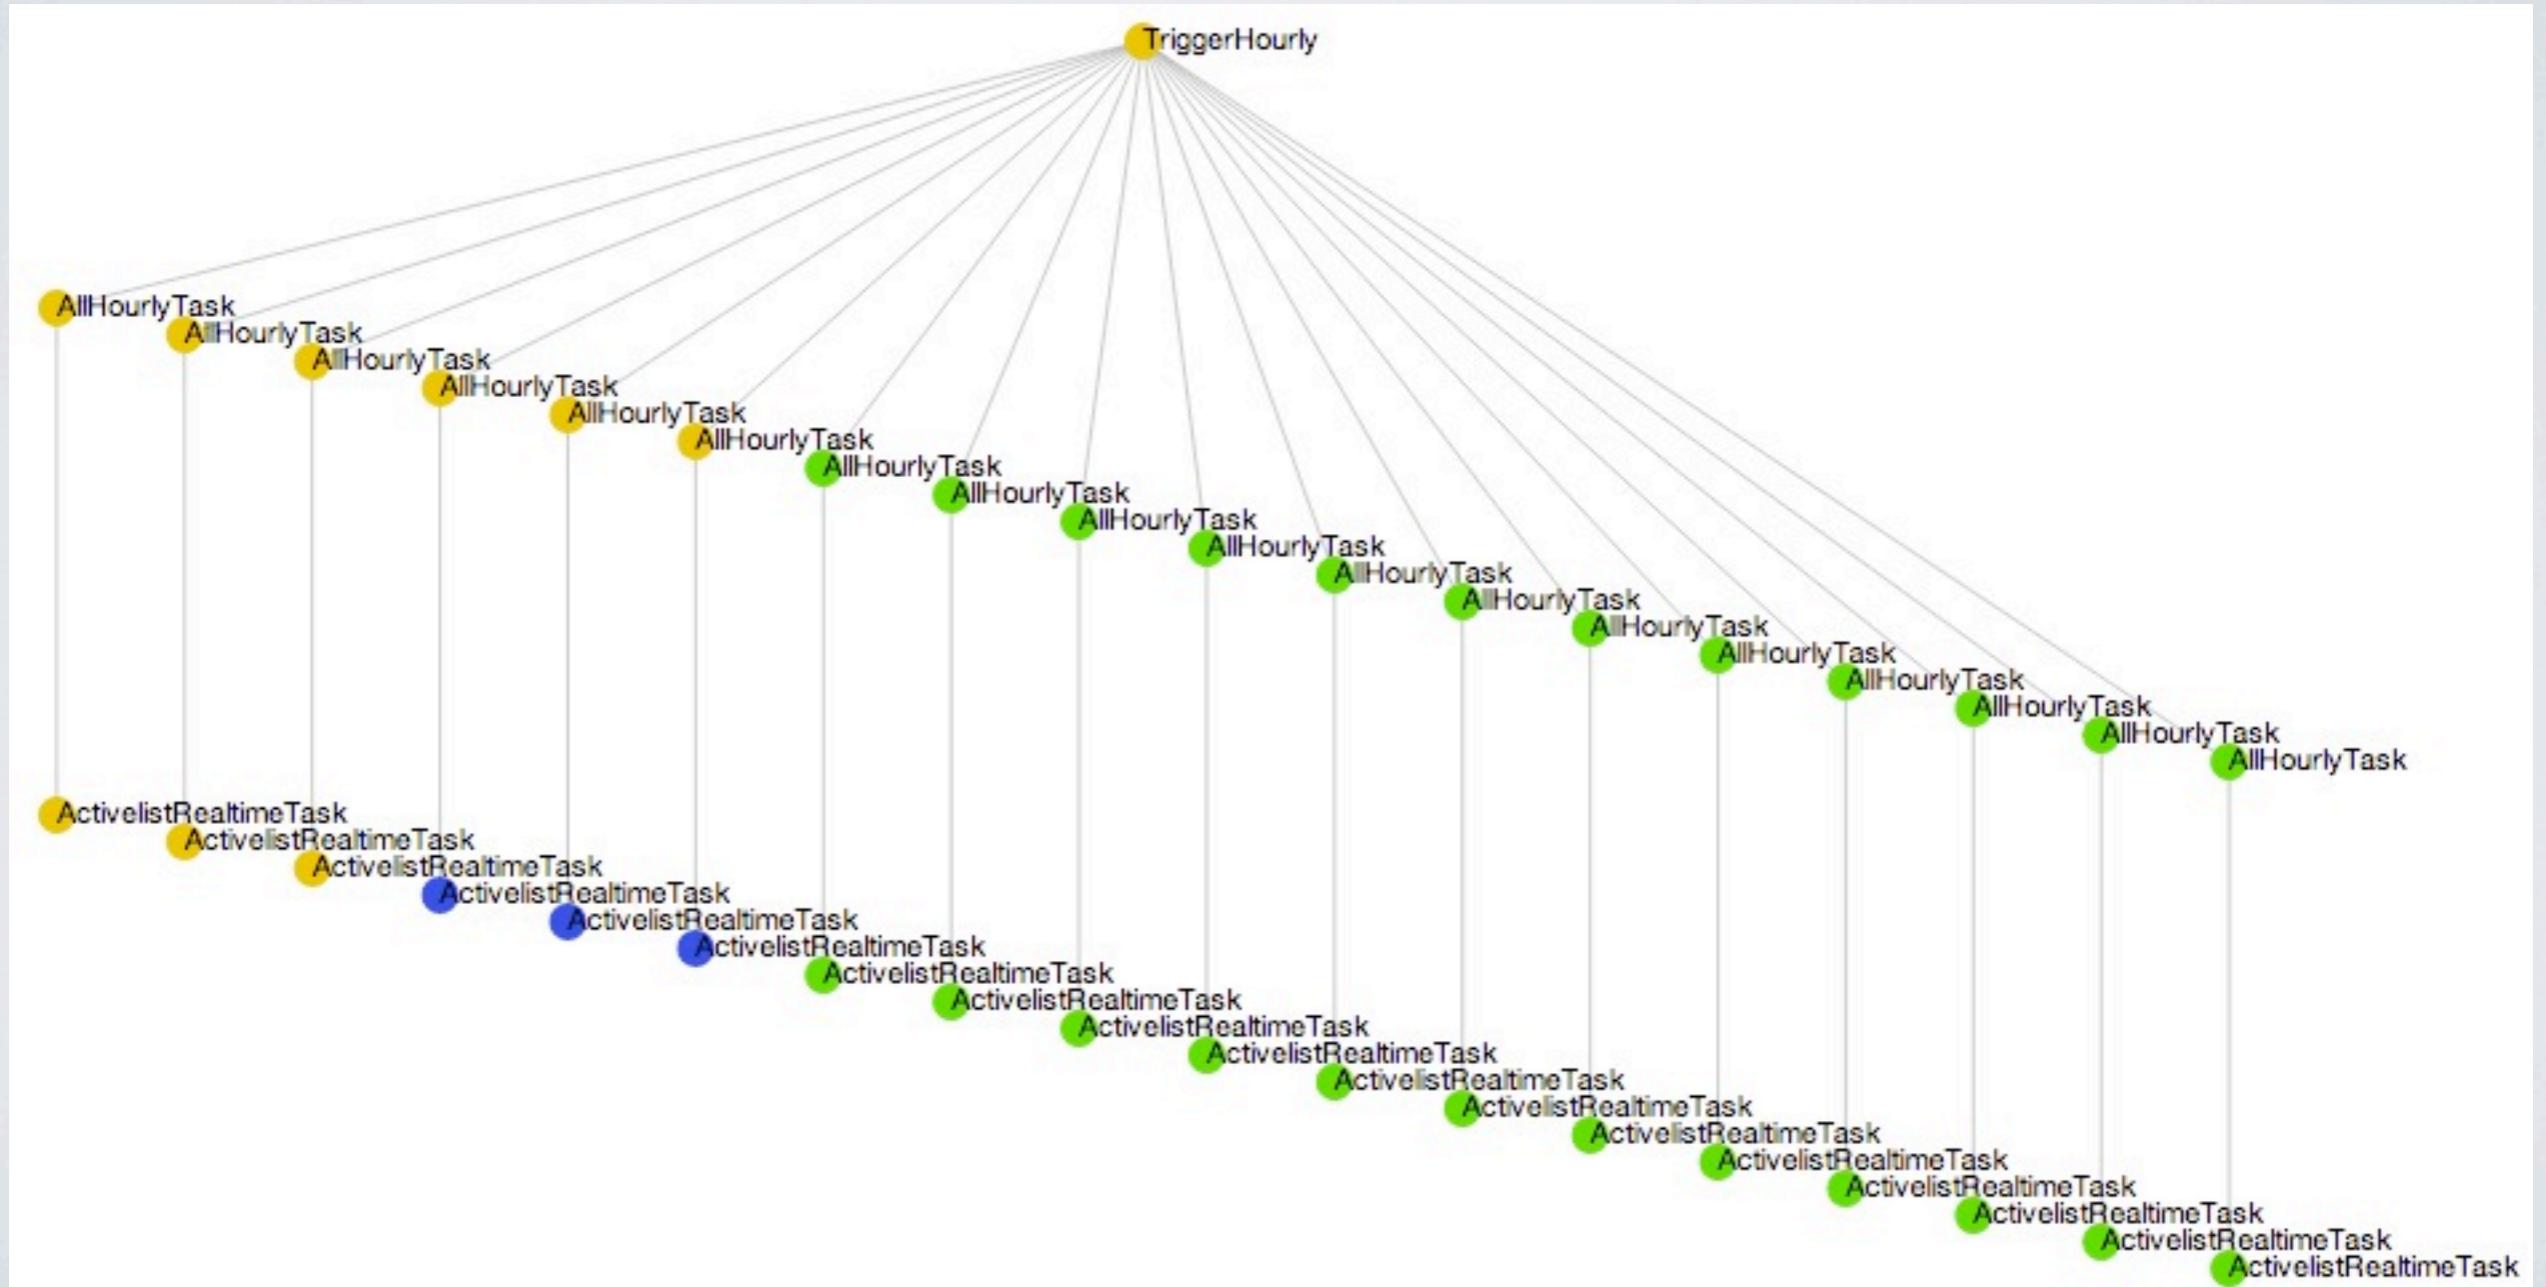


MOAR SCREENSHOTS

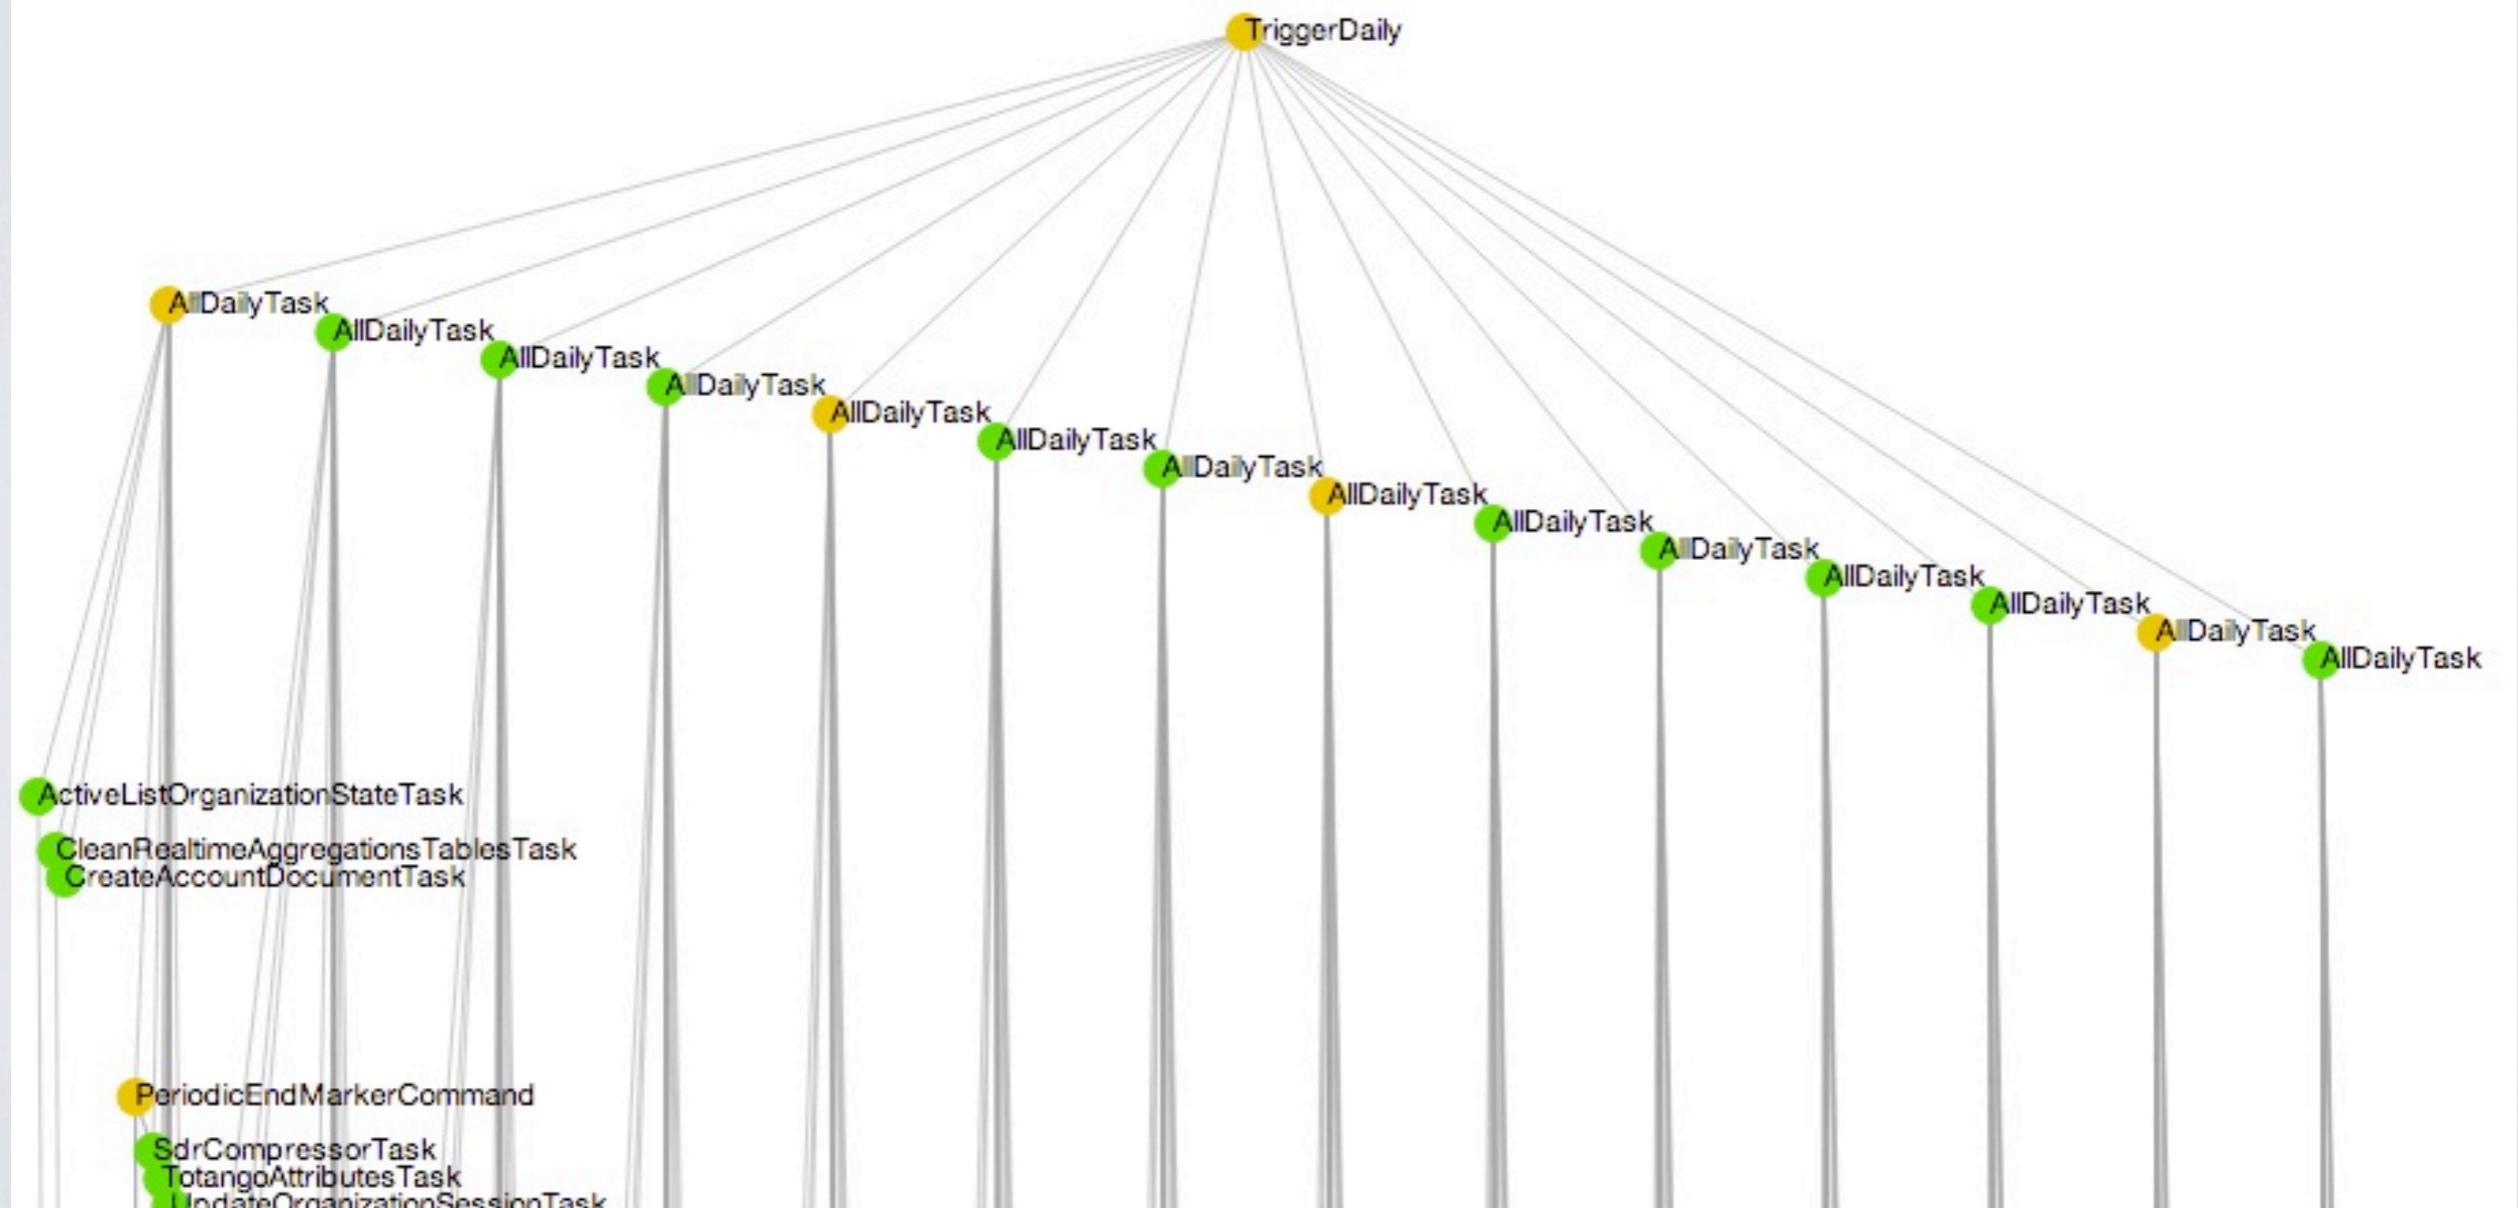


MOAR

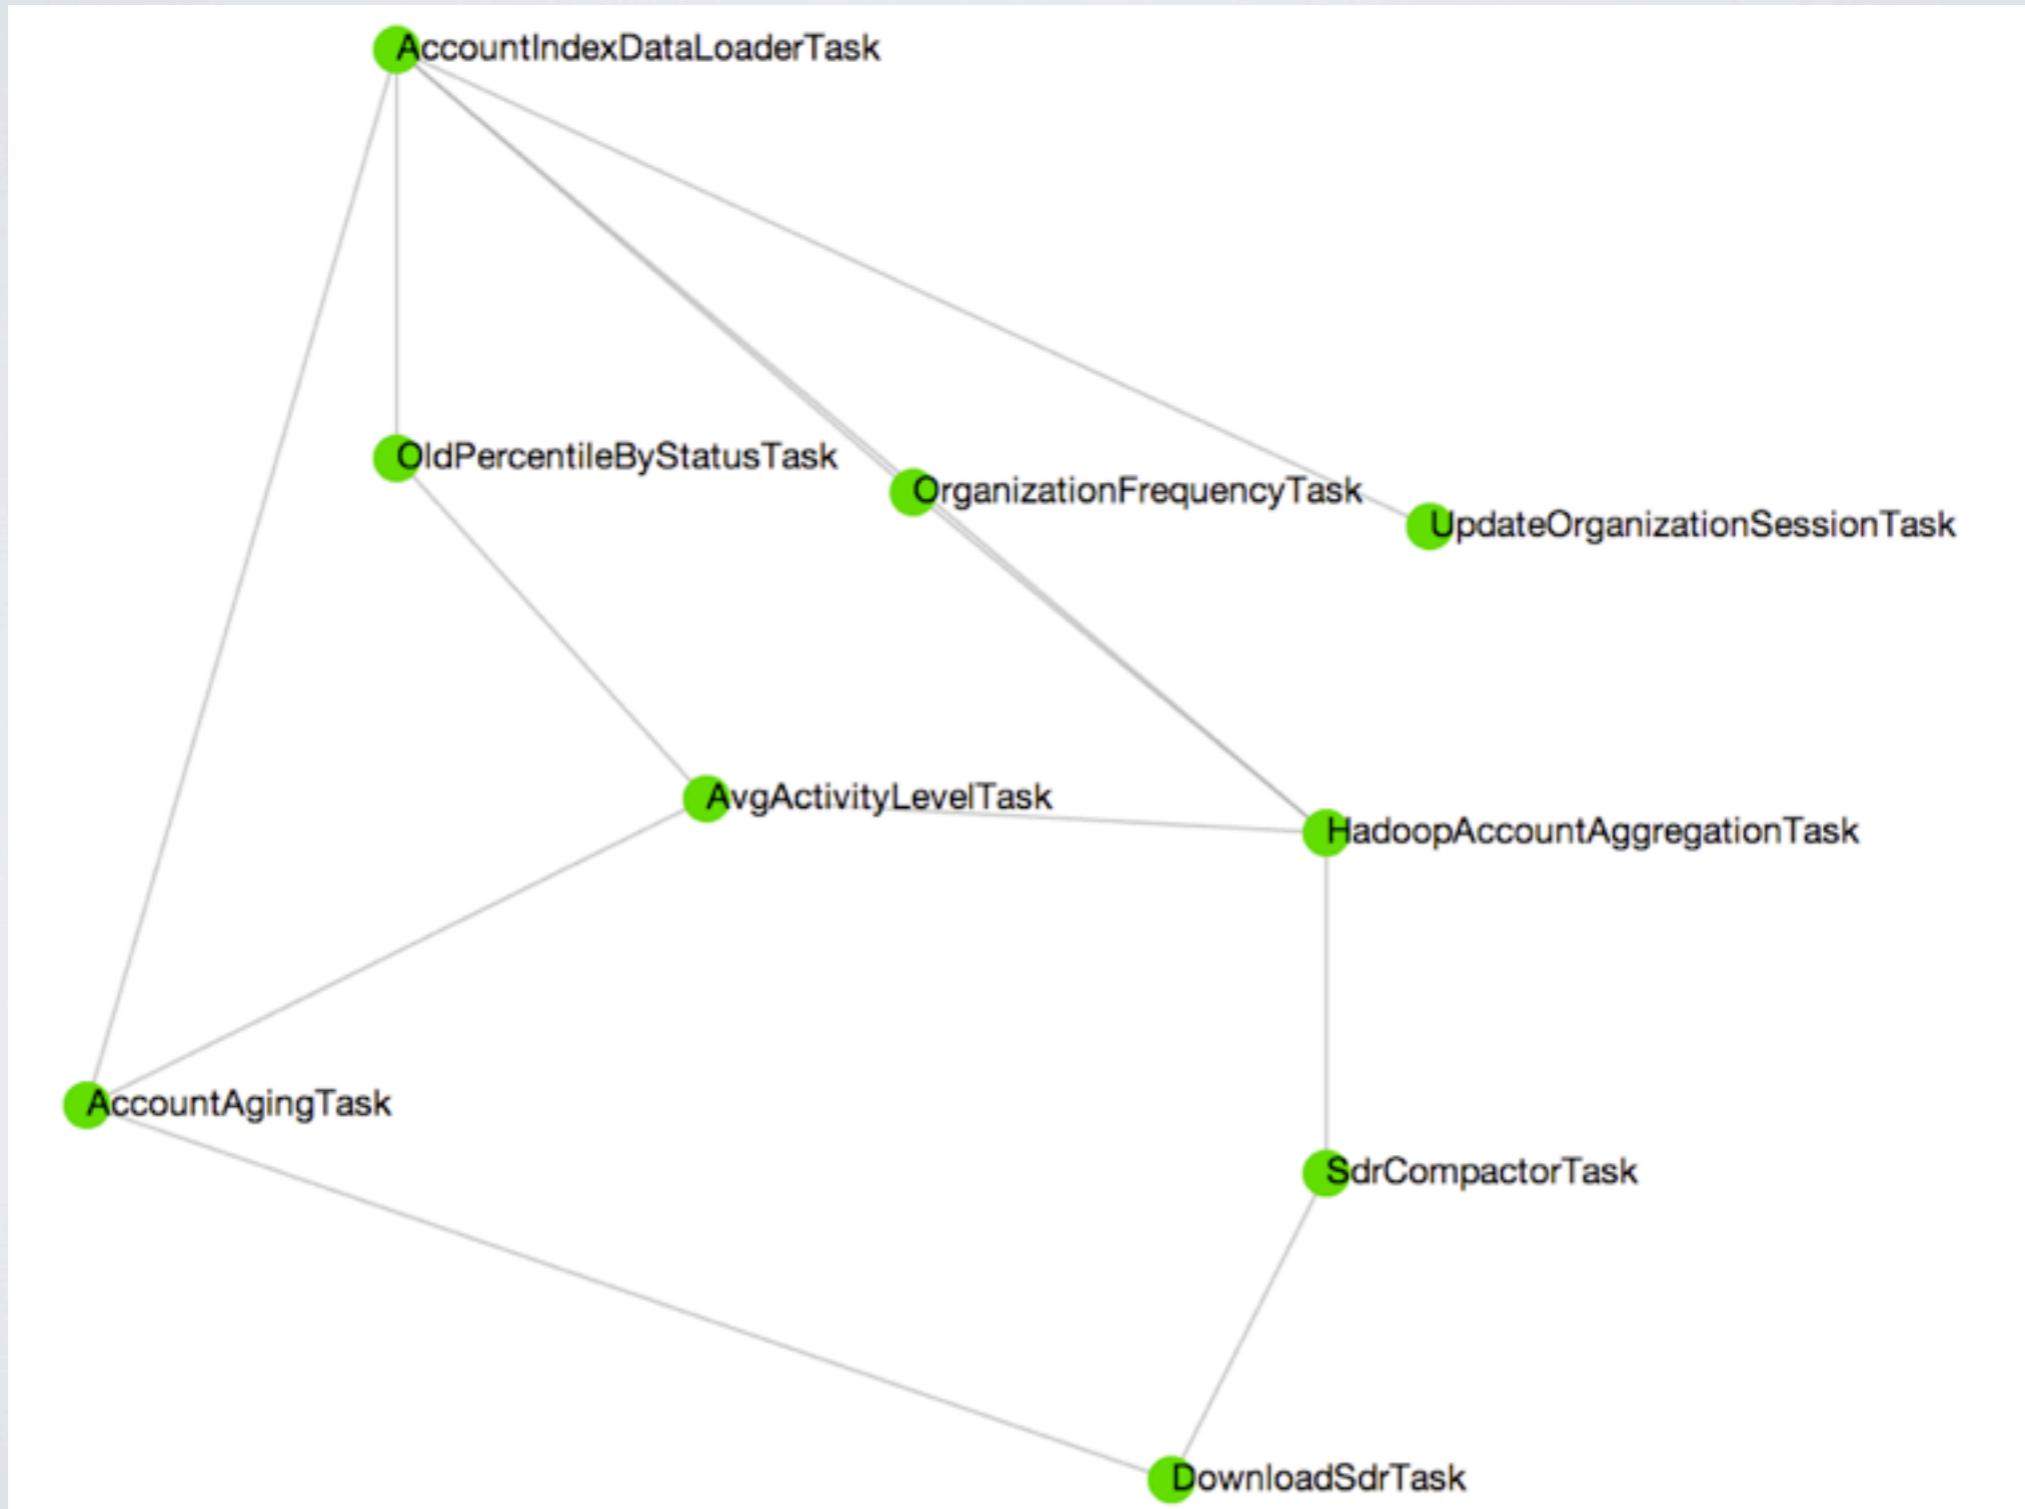
MOAR SCREENSHOTS



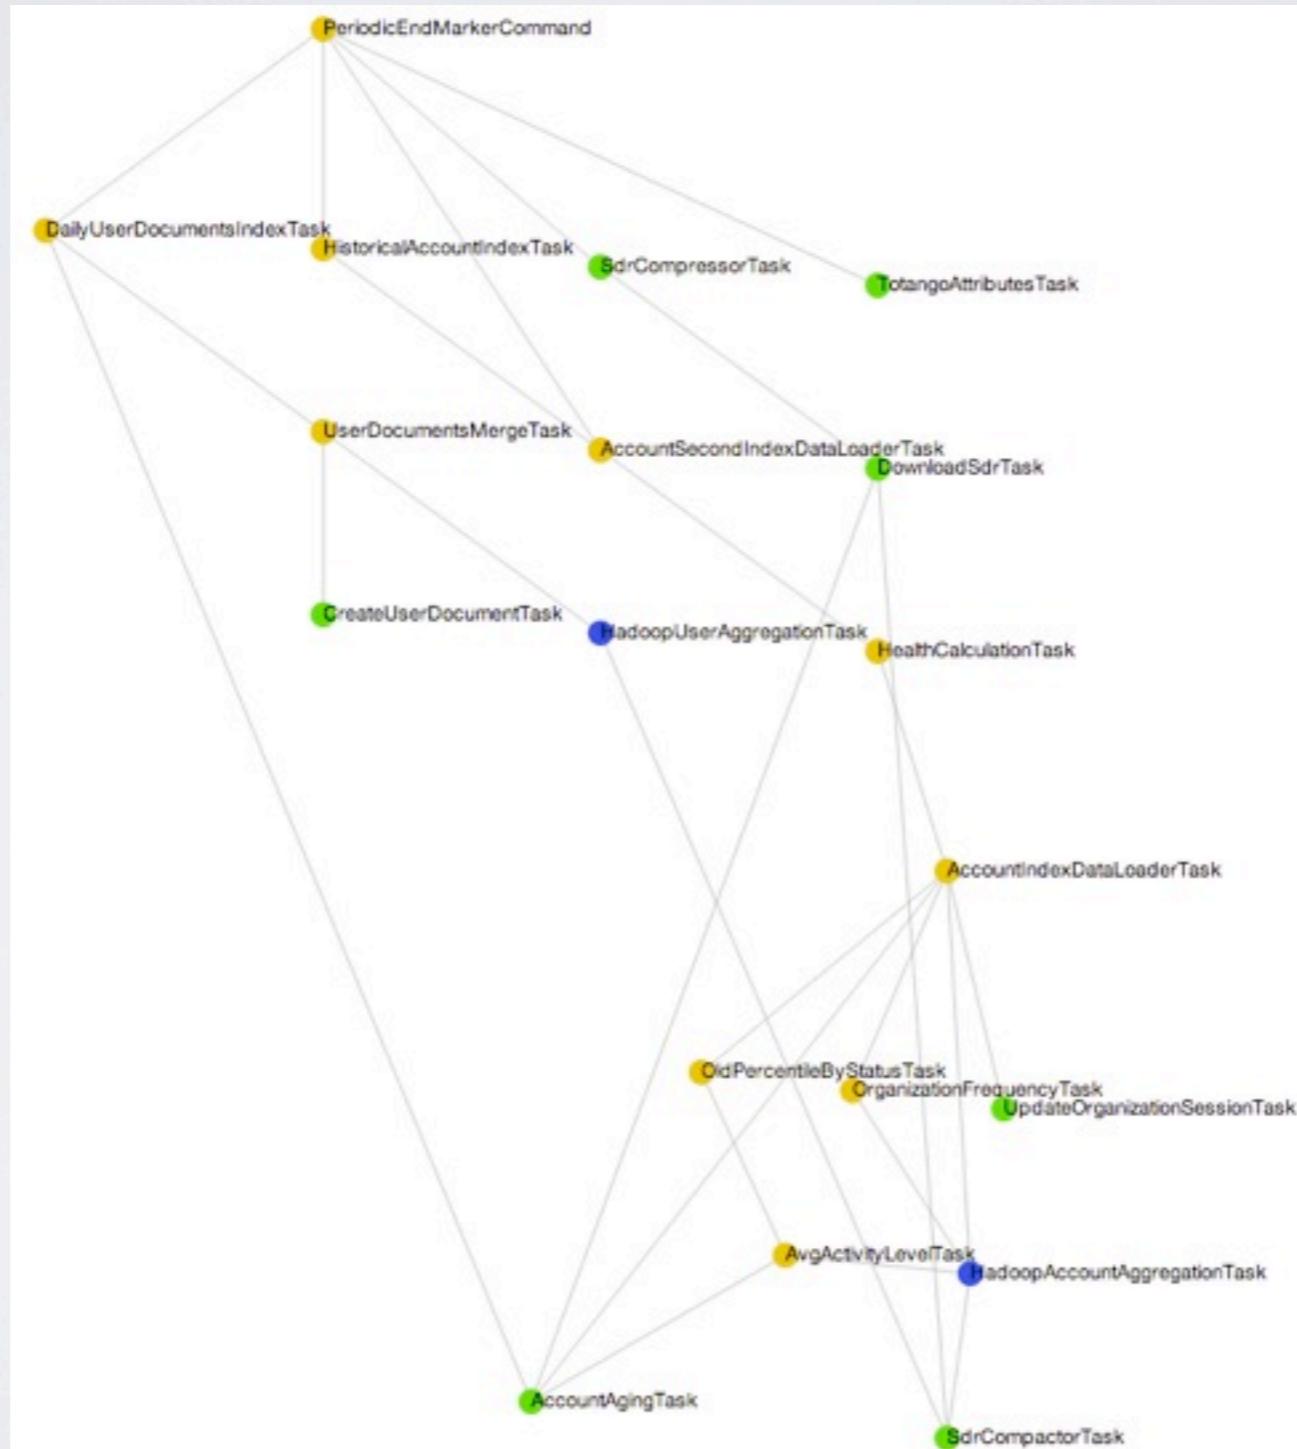
MOAR SCREENSHOTS



MOAR SCREENSHOTS



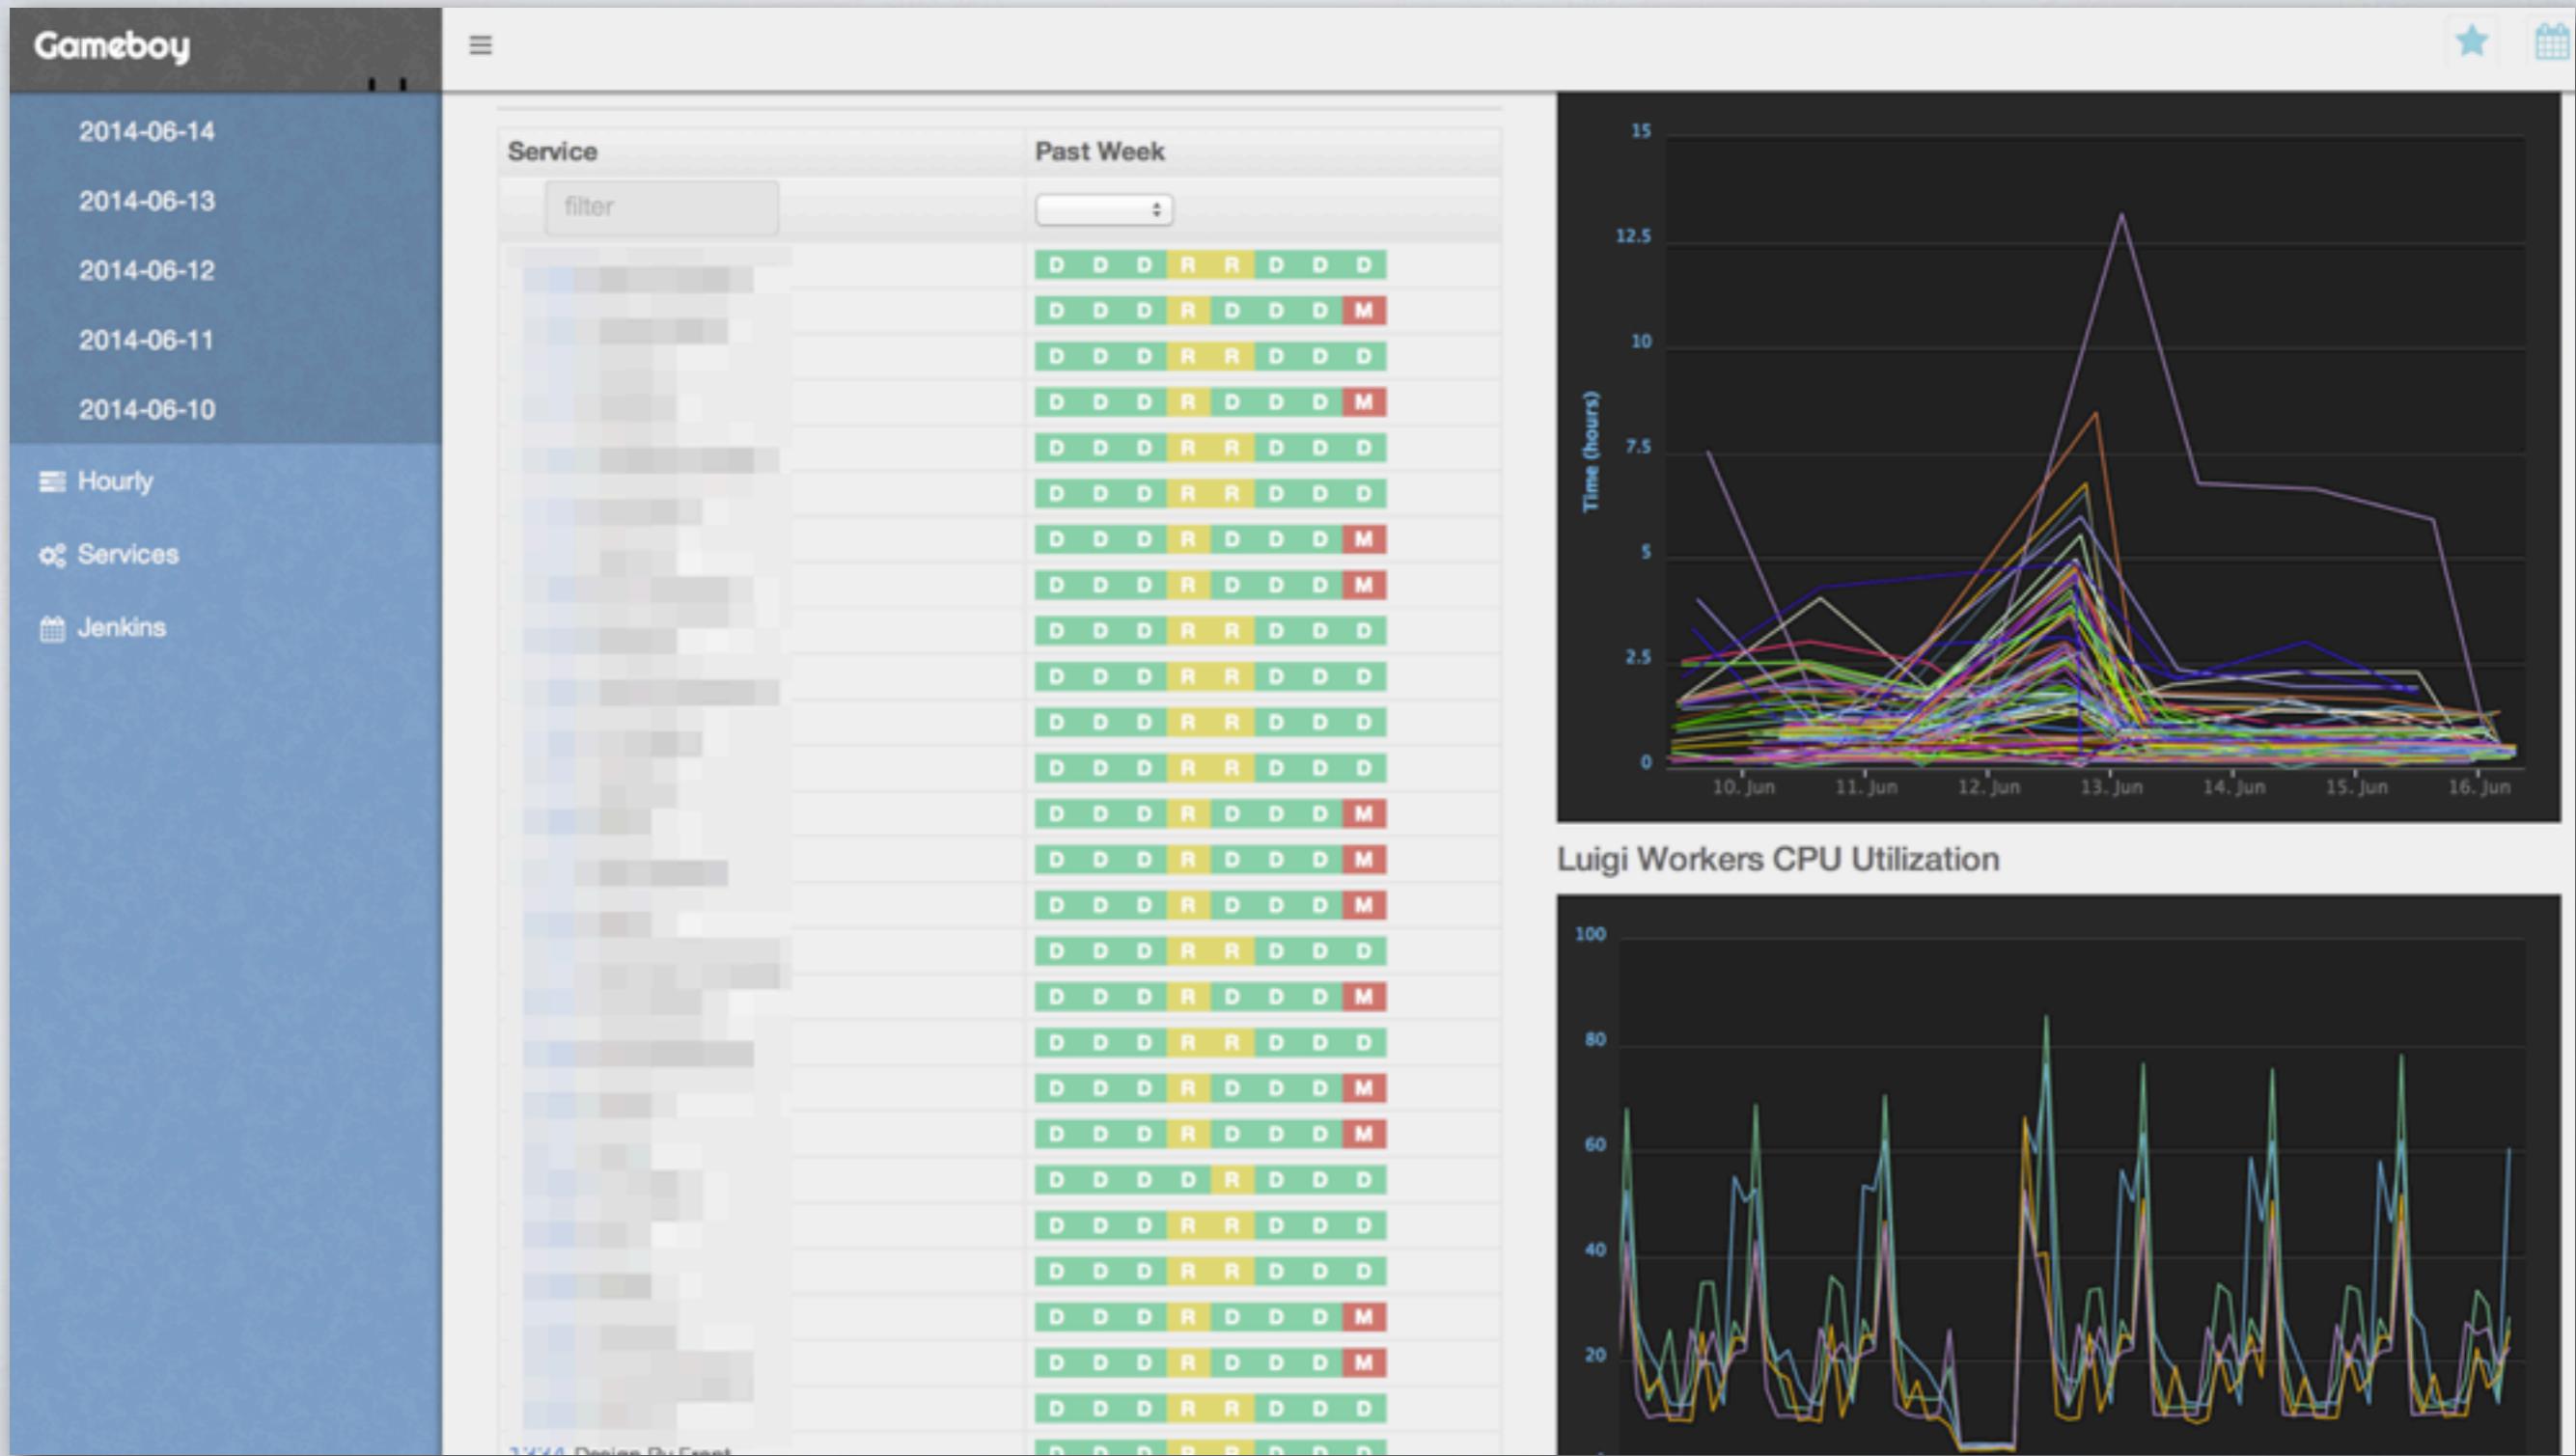
MOAR SCREENSHOTS



GAMEBOY!!!



AND... GAMEBOY



GAMEBOY

Gameboy

Home / Dashboard

TODAY

130 (30 pending/running)	2351 (56 pending/running)	0 FAILED
DAILY	HOURLY	
96 services (96 in luigi)	96 services (96 in luigi)	

JENKINS QUEUE 0

PAST WEEK DAILIES

Service	Past Week
	D D D D D D D P
	D D D D D D D P
	D D D D D D D P
M	D D D D D D D D
M	D D D D D D D D
D	D D D D D D D P
D	D D D D D D D P
M	D D D D D D D D
D	D D D D D D D P
M	D D D D D D D D
M M M M D D D D	
M	D D D D D D D D
M	D D D D D D D D
M	D D D D D D D P
D	D D D D D D D P

RUNNING/PENDING/FAILED TODAY

Task	Status	Service
DailyUserDocumentsIndexTask {"service_id": "1", "force_run": "False", "day": "2014-06-22"}	PENDING	1
PeriodicEndMarkerCommand {"service_id": "1", "force_run": "False", "day": "2014-06-22"}	PENDING	1
AllDailyTask {"service_id": "1", "force_run": "False", "day": "2014-06-22", "tz_string": "-0700"}	PENDING	1
HistoricalAccountIndexTask {"service_id": "1", "force_run": "False", "day": "2014-06-22"}	PENDING	1
AccountSecondIndexDataLoaderTask {"service_id": "1", "force_run": "False", "day": "2014-06-22"}	PENDING	1
HealthCalculationTask {"service_id": "1", "force_run": "False", "day": "2014-06-22"}	PENDING	1
AccountIndexDataLoaderTask		

Gameboy Status

- 200 admin_db
- 200 history_db
- 200 jenkins
- 200 luigi
- 200 maint_db
- 200 status

Jenkins Status

- enabled daily-cron
- enabled daily-manual
- enabled hourly-cron

Luigi Workers CPU Utilization

GAMEBOY

Gameboy

Dashboard

Daily

All Dailies

2014-06-23

2014-06-22

2014-06-21

2014-06-20

2014-06-19

2014-06-18

2014-06-17

Hourly

Services

Jenkins

Home / Daily / Daily on 2014-06-21

DAILY FOR 2014-06-21

ALL JOBS

Service	Start	Last update	Duration	Status
	22nd - 07:01:56	22nd - 09:23:13	2 hours	DONE
	22nd - 07:02:05	22nd - 10:56:59	4 hours	DONE
	22nd - 07:01:41	22nd - 12:56:49	6 hours	DONE
	22nd - 12:31:40	22nd - 02:33:12	2 hours	DONE
	22nd - 04:01:45	22nd - 04:42:14	40 minutes	DONE
	22nd - 07:01:56	22nd - 09:19:24	2 hours	DONE
	22nd - 07:02:03	22nd - 10:06:33	3 hours	DONE
	22nd - 12:31:39	22nd - 01:14:05	42 minutes	DONE
	22nd - 07:01:54	22nd - 08:41:19	2 hours	DONE
	21st - 09:31:46	21st - 10:23:40	an hour	DONE
	21st - 09:31:44	21st - 10:25:55	an hour	DONE
	22nd - 04:01:42	22nd - 04:55:17	an hour	DONE
	21st - 10:31:40	21st - 11:35:10	an hour	DONE
	22nd - 07:02:01	22nd - 09:40:05	3 hours	DONE
	22nd - 06:01:41	22nd - 02:48:28	9 hours	DONE

GAMEBOY IS

- A Totango specific controller for Luigi
 - The transition process (to Luigi)
 - Provide **high level** overview
 - Manual **re-run** of tasks
 - Monitor **progress, performance, run times, queues, worker load** etc...
 - Implemented using **Flask** and **AngularJS**



IS GAMEBOY OPEN SOURCE

- Well, no. At least not right away
- Right now it's very totango-specific.
 - Integrations to Librato-metrics
 - Queries on Totango Databases
 - Uses Jenkins for controlling executions
 - Displays Totango's Account metadata (totango's business logic)
- Maybe some other day...

WHAT ELSE IS OUT THERE?



WHAT ELSE IS OUT THERE?

- Oozie
- Azkaban
- AWS Data Pipeline
- Chronos
- spring-batch
- Dataswarm (facebook)
- River (outbrain internal)
- What's your favorite WF engine? (did you build one?)

MY OTHER PROJECTS

- <https://github.com/hector-client/hector>
- <https://github.com/rantav/flask-restful-swagger>
- <https://github.com/sebastien/monitoring>
- <https://github.com/rantav/meteor-migrations>
- <https://github.com/rantav/node-github-list-packages>
- <https://github.com/rantav/devdev>



REFS

- <https://github.com/spotify/luigi>
- Facebook's Dataswarm <https://www.youtube.com/watch?v=M0VCbhfQ3HQ>
- Outbrain's River <https://www.youtube.com/watch?v=EzsckTggDiM>